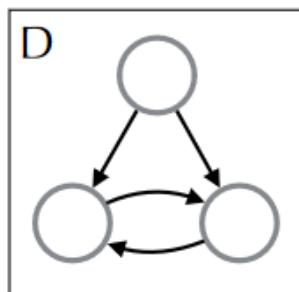





Pre-class Puzzle

- For each of these Control Flow Graphs (CFGs), what is a C program that corresponds to it?



Take some time to try to figure these out---I'll ask for volunteers to share their answers around 10:05

CS443: Compiler Construction

Lecture 15: Loop Optimization

Stefan Muller

Based on material by Steve Zdancewic, Stephen Chong and Greg Morrisett

Loop optimizations are especially important!

- Programs spend most of the time in loops
- Lots of loop optimizations:
 - Loop invariant removal
 - Induction variable elimination
 - Loop unrolling
 - Loop fusion
 - Loop fission
 - Loop peeling
 - Loop interchange
 - Loop tiling
 - Loop parallelization
 - Software pipelining

Invariant removal: Don't recompute things in a loop

```
10:  
    %i = bitcast i32 0 to i32  
    br %l1  
  
l1:  
    %i = add i32 %i 1  
    %t = add i32 %a %b  
    %el = getelementptr i32, i32* %arr, i32 %i  
    store i32 %t, i32* %el  
    %lt = icmp lt i32 %i %N  
    br i1 %lt, label %l1, label %l2  
  
l2:  
    ret %t
```

Invariant removal: Don't recompute things in a loop

```
l0:  
    %i = bitcast i32 0 to i32  
    %t = add i32 %a %b  
    br %l1  
  
l1:  
    %i = add i32 %i 1  
    %el = getelementptr i32, i32* %arr, i32 %i  
    store i32 %t, i32* %el  
    %lt = icmp lt i32 %i %N  
    br i1 %lt, label %l1, label %l2  
  
l2:  
    ret %t
```

Loop induction variable: Avoid recompilation based on loop induction variables

```
10:  
    %i = bitcast i32 0 to i32  
11:  
    %t1 = mul i32 %i 4  
    %t2 = add i32 %a %t1  
    %s = add i32 %s %t2  
    %i = add i32 %i 1  
    %lt = icmp lt %i 100  
    br i1 %lt, label %l1, label %l2  
12: ...
```

t1 is always equal to $i * 4$

Loop induction variable: Avoid recomputation based on loop induction variables

10:

```
%i = bitcast i32 0 to i32  
%t1 = bitcast i32 -4 to i32
```

t2 is always $a + i * 4$

11:

```
%t1 = add i32 %t1 4  
%t2 = add i32 %a %t1  
%s = add i32 %s %t2  
%i = add i32 %i 1  
%lt = icmp lt %i 100  
br i1 %lt, label %l1, label %l2
```

12: ...

Loop induction variable: Avoid recompilation based on loop induction variables

10:

```
%i = bitcast i32 0 to i32
%t1 = bitcast i32 -4 to i32
%t2 = bitcast i32 %a to i32
```

Can eliminate t1!

11:

```
%t1 = add i32 %t1 4
%t2 = add i32 %t2 4
%s = add i32 %s %t2
%i = add i32 %i 1
%lt = icmp lt %i 100
br i1 %lt, label %l1, label %l2
```

12: ...

Loop induction variable: Avoid recomputation based on loop induction variables

10:

```
%i = bitcast i32 0 to i32  
%t2 = bitcast i32 %a to i32
```

Can eliminate i!

11: %t2 = add i32 %t2 %4

```
%s = add i32 %s %t2
```

```
%i = add i32 %i 1
```

```
%lt = icmp lt %i 100
```

```
br i1 %lt, label %l1, label %l2
```

12: ...

Loop induction variable: Avoid recompilation based on loop induction variables

10:

```
%i = bitcast i32 0 to i32
%t2 = bitcast i32 %a to i32
%endt2 = add i32 %a 400
```


11:

```
%t2 = add i32 %t2 %4
%s = add i32 %s %t2
%lt = icmp lt %t2 %endt2
br i1 %lt, label %l1, label %l2
```


12: ...

Before we can optimize loops, we have to find them


- In C (without goto): easy!
- In LLVM: surprisingly hard!
- **Definition (loop):**
 - Subset S of nodes in CFG
 - Designated “header” node h
 - S is strongly connected
 - No edge from outside S to $S \setminus \{h\}$



Loops can be nested




Non-example: (there can't be a header)






- (But this can't arise in C/C++ without goto)

A node d dominates n if every path (from start) to n must go through d




We can define loops based on dominators

- A **back edge** is an edge from n to a dominator d
- If there's a back edge $n \rightarrow d$, there is a **loop** consisting of the set of nodes x such that d dominates x and there is a path from x to n not including d



Example

- For each node, what nodes does it dominate? Back edges?
- a dominates a, b, c, d, e, f, g, h
- b dominates b, c, d, e, f, g, h
- c dominates c, e
- d dominates d
- f dominates f, g, h
- g dominates g, h
- h dominates h,
- Back edges: g \rightarrow b, h \rightarrow a




We can calculate dominators with a dataflow analysis!

- $\text{out}[n] = \text{set of nodes that dominate } n$
 - $\text{in}[n] := \bigcap_{n' \in \text{pred}[n]} \text{out}[n']$
 - $\text{out}[n] := \text{in}[n] \cup \{n\}$
- Forward must analysis: initialize $\text{out}[n]$, $\text{in}[n]$ to all nodes

We can represent dominators with a “dominator tree”

- Edge to n from its “immediate dominator” (dominator other than n that is dominated by other dominators other than n)



Identifying loop invariants

- An instruction $\%x = \text{opc } op_1, op_2, \dots, op_N$ represented by a node n is invariant for a loop if for each operand op_i :
 - op_i is constant, or
 - all definitions of op_i that reach n are outside the loop, or
 - only one definition reaches op_i and it is a loop invariant

Loop invariant example from before

```
l0:  
    %i = bitcast i32 0 to i32  
    br %l1  
  
l1:  
    %i = add i32 %i 1  
    %t = add i32 %a %b  
    %el = getelementptr i32, i32* %arr, i32 %i  
    store i32 %t, i32* %el  
    %lt = icmp lt i32 %i %N  
    br i1 %lt, label %l1, label %l2  
  
l2:  
    ret %t
```

Actually moving (*hoisting*) invariants out of the loop is pretty tricky

- Move to a “pre-header” (CFG node before header)

Need to make sure hoisting wouldn't interfere with other uses!

```
10:  
    %i = bitcast i32 0 to i32  
    br %l1  
  
11:  
    %i = add i32 %i 1  
    %el = getelementptr i32, i32* %arr, i32 %i  
    store i32 %t, i32* %el  
    %t = add i32 %a %b  
    %lt = icmp lt i32 %i %N  
    br i1 %lt, label %l1, label %l2  
  
12:  
    ret %t
```

Need to make sure hoisting wouldn't interfere with other uses!

- $n := \%x = \text{opc op1, op2, ..., opN}$ is safe to hoist if:
 - n dominates all loop exits at which $\%x$ is live, and
 - there is only one definition of x in the loop, and
 - x is not live at the pre-header

```
for (i = 0; i < 100; i += 2) {          t = 0;
    t = a + b;                      while(1) {
    a[i] = t;                      break;
    t = a - b;                      t = a + b;
    a[i + 1] = t;                  a[i] = t;
}                                            }
```

Multiple definitions of t!

t doesn't dominate
the break

Loop Unrolling: Copy over the body of a loop

```
for (int i = 0; i < n; i++) {  
    a[i] = i;  
}
```



```
//Handle the first few in case n not a multiple of 3  
for (int i = 0; i < n % 3; ++i) a[i] = i;  
for (; i < n; i+=3) {  
    a[i] = i;  
    a[i + 1] = i + 1;  
    a[i + 2] = i + 2;  
}
```


Why is this an optimization?

Loop unrolling: costs and benefits

- Benefits:
 - Amortize tests, jumps over more instructions
- Costs:
 - Program size increases (why is this a problem?)

Loop Peeling: “Peel off” the first or last N iterations of a loop

```
for (int i = 0; i < N; i++) {  
    if (i <= 1) {  
        a[i] = i;  
    } else {  
        a[i] = a[i - 1] + a[i - 2];  
    }  
}
```



```
a[0] = 0;  
a[1] = 1;  
for (int i = 2; i < N; i++) {  
    a[i] = a[i - 1] + a[i - 2];  
}
```

Loop Interchange: Swap order of nested loops

```
for (int i = 0; i < w; i++) {  
    for (int j = 0; j < h; j++) {  
        sum += a[j][i];  
    }  
}
```



```
for (int j = 0; j < h; j++) {  
    for (int i = 0; i < w; i++) {  
        sum += a[j][i];  
    }  
}
```

Why is this an optimization?