



Feel free to take candy!
(Subject to the following restrictions)



- For every pair of people, if your **first or last** names start with the same letter, you can't take the same kind of candy.
- Stefan has already taken a Kit Kat

(Don't worry, if this only leaves you with candy you don't like/are allergic to/etc., you can get more)

CS443: Compiler Construction

Lecture 19: Register Allocation

Stefan Muller

Based on material by Steve Zdancewic

Register allocation: going from unlimited temporaries to fixed number of registers

Register	ABI Name	
x0	zero	
x1	ra	
x2	sp	
x3	gp	
x4	tp	
x5-7	t0-2	
x8	s0/fp	
x9	s1	
x10-11	a0-1	
x12-17	a2-7	
x18-27	s2-11	
x28-31	t3-t6	

The table shows the mapping between registers and their ABI names. The registers are color-coded into four groups:

- Special purpose:** x0, x1, x2, x3, x4 (blue)
- General purpose:** x5-7, x8, x9 (light blue)
- Sometimes special purpose (by convention):** x10-11, x12-17 (light green)
- General purpose:** x18-27, x28-31 (medium blue)

Brackets on the right side of the table group these categories.

Find: mapping from program variables to registers

- What if there aren't enough registers?

```
int annoying(int[] a) {  
    int v0 = a[0];  
    int v1 = a[1];  
    int v2 = a[2];  
    int v3 = a[3];  
    int v4 = a[4];  
    int v5 = a[5];  
    int v6 = a[6];  
    int v7 = a[7];  
    int v8 = a[8];  
    int v9 = a[9];  
  
    ...  
    return v0 + v1 + v2 + v3 + v4 + ...  
}
```

Find: mapping from program variables to
(registers \cup stack locations)

```
type alloc_res = InReg of R.reg
  | OnStack of int (* stack slot, 0-N *)
  | InMem of R.symbol (* globals on heap *)
```

“spill”

Many quality metrics for allocation

- Program semantics is preserved (i.e. the behavior is the same)
- Register usage is maximized
- Moves between registers are minimized
- Calling conventions / architecture requirements are obeyed

Recall: A variable is “live” when its value is needed

```
int f(int x) {  
    int a = x + 2;           ← x is live  
    int b = a * a;          ← a and x are live  
    int c = b + x;          ← b and x are live  
    return c;                ← c is live  
}
```

Liveness analysis is based on uses and definitions

- For a node/statement s define:
 - $\text{use}[s]$: set of variables used (i.e. read) by s
 - $\text{def}[s]$: set of variables defined (i.e. written) by s
- Examples:
 - $a = b + c$ $\text{use}[s] = \{b,c\}$ $\text{def}[s] = \{a\}$
 - $a = a + 1$ $\text{use}[s] = \{a\}$ $\text{def}[s] = \{a\}$

Liveness analysis as a dataflow analysis (Steps 1-2)

- Facts: Live variables
- $\text{gen}[n] = \text{use}[n]$
- $\text{kill}[n] = \text{def}[n]$
- Constraints:
 - $\text{in}[n] \supseteq \text{gen}[n]$
 - $\text{out}[n] \supseteq \text{in}[n']$ if $n' \in \text{succ}[n]$
 - $\text{in}[n] \supseteq \text{out}[n] / \text{kill}[n]$

Liveness analysis as a dataflow analysis (Steps 3-4)

- Equations:
 - $\text{out}[n] := \bigcup_{n' \in \text{succ}[n]} \text{in}[n']$
 - $\text{in}[n] := \text{gen}[n] \cup (\text{out}[n] / \text{kill}[n])$
- Initial values:
 - $\text{out}[n] := \emptyset$
 - $\text{in}[n] := \emptyset$

For register allocation: $\text{live}(x)$

- $\text{live}(x) = \text{set of variables that are live-in to the definition of } x$
 - (assuming SSA)

Linear Scan: a simple, greedy algorithm

1. Compute liveness information: `live(x)`
2. Let `regs` be the set of usable registers
3. Maintain "layout" `alloc` that maps uids to `alloc_reg`
4. Scan through the program. For each instruction that defines a var `x`
 - `used = {r | reg r = alloc(y) s.t. y ∈ live(x)}`
 - `available = regs - used`
 - If `available` is empty: *// no registers available, spill*
`alloc(x) := OnStack n; n := !n + 1`
 - Otherwise, pick `r` in `available`: *// choose an available register*
`alloc(x) := InReg r`

Linear Scan Example (registers: r0, r1, r2)


```
int f(int x) {          Available
    int a = x + 2;      r0, r1, r2      a -> r0
    int b = a * a;      r1, r2      b -> r1
    int c = b + a;      r2      c -> r2
    return c;
}
```

Linear scan is OK, but we can do better

Who had “reduce it to a graph problem” on their CS Bingo card?

- Nodes of the graph are variables
- Edges connect variables that *interfere* with each other
 - Two variables interfere if their live ranges intersect (i.e. there is an edge in the control-flow graph across which they are both live).
- Register assignment is a *graph coloring*.
 - A graph coloring assigns each node in the graph a color (register)
 - Any two nodes connected by an edge must have different colors.
- Example:

```
%b1 = add i32 %a, 2  
%c = mult i32 %b1, %b1  
%b2 = add i32 %c, 1  
%ans = add i32 %b2, %a  
return %ans;
```



Interference Graph

2-Coloring of the graph
red = r8
yellow = r9

Heuristics for graph coloring come down to order in which you color nodes

- Linear Scan: Order of definitions in program
- *Simplification*: (Roughly) color high degree nodes first

Coloring by simplification

1. **Build** Interference Graph
2. **Simplify** the graph by removing nodes one at a time, putting them on a stack
3. **Select** colors for nodes in order of the stack

We don't want to treat move instructions as conflicts/interference

```
%a = inttoptr i32* %aptr to i32
%b = add i32 %a 8
%bptr = ptrtoint i32 %b to i32*
%c = load i32, i32* %aptr
%d = load i32, i32* %bptr
```

%a and %aptr are live at the same time, but can (and should) be in the same register

Steps for a simple graph-coloring allocator

1. **Build** interference graph
2. **Simplify** graph by removing nodes one at a time, pushing them on a stack, until all nodes are on stack
3. As we simplify, identify nodes to potentially **spill**
4. **Select** colors/registers for nodes (in reverse order they were pushed to the stack)

Build interference graph

- For each instruction:
 - If the inst defines a variable a , with b_1, \dots, b_n live-out:
 - If the instruction is not a move, add edges $(a, b_1), \dots, (a, b_n)$
 - If the instruction is a move $a = c$, add edges $\{(a, b_i) \mid b_i \neq c\}$

Coloring by simplification: Simplify/Spill


- Let K = number of registers
- Let S = empty stack
- While graph not empty:
 - If there exists a node m with fewer than K neighbors:
 - Remove m from the graph, push it on S
 - Guaranteed that we will be able to find a color for m
 - Otherwise:
 - Pick a node m , remove it from the graph, push it on S (we may end up spilling it)

Coloring by simplification: Select

- While S not empty:
 - Pop m from S
 - If there is a color (register) available for m :
 - Choose an available color (register) for m and add it back to the graph
 - Otherwise:
 - Spill m – put it in the next stack slot


Graph Coloring Example (Appel)

```
g = mem[j + 12]
h = k - 1
f = g * h
e = mem[j + 8]
m = mem[j + 16]
b = mem[f]
c = e + 8
d = c          % Move
k = m + 4
j = b          % Move
                % d, k, j live-out
```




Graph Coloring Example (Appel)

g




Graph Coloring Example (Appel)

h
g




Graph Coloring Example (Appel)

k
h
g




Graph Coloring Example (Appel)

d
k
h
g




Graph Coloring Example (Appel)

j
d
k
h
g




Graph Coloring Example (Appel)

e
j
d
k
h
g




Graph Coloring Example (Appel)

f
e
j
d
k
h
g



Graph Coloring Example (Appel)

b
f
e
j
d
k
h
g



Graph Coloring Example (Appel)


c
b
f
e
j
d
k
h
g



m


Graph Coloring Example (Appel)

m
c
b
f
e
j
d
k
h
g




Graph Coloring Example (Appel)

c
b
f
e
j
d
k
h
g




Graph Coloring Example (Appel)

b
f
e
j
d
k
h
g




Graph Coloring Example (Appel)

f
e
j
d
k
h
g




Graph Coloring Example (Appel)

e
j
d
k
h
g




Graph Coloring Example (Appel)

j
d
k
h
g




Graph Coloring Example (Appel)

d
k
h
g




Graph Coloring Example (Appel)

k
h
g




Graph Coloring Example (Appel)

h
g




Graph Coloring Example (Appel)

g




Graph Coloring Example (Appel)



Graph Coloring Example (Appel)


```
g = mem[j + 12]
h = k - 1
f = g * h
e = mem[j + 8]
m = mem[j + 16]
b = mem[f]
c = e + 8
d = c
k = m + 4
j = b
```



Graph Coloring Example (Appel)


```
r4 = mem[r3 + 12]
r2 = r1 - 1
r2 = r4 * r2
r4 = mem[r3 + 8]
r1 = mem[r3 + 16]
r2 = mem[r2]
r3 = e + 8
r4 = r3 ←
r1 = r1 + 4
r3 = r2 ←
```

Next time:
Avoid these




Graph Coloring Example (Appel)

c
h




Graph Coloring Example (Appel)

g
c
h




Graph Coloring Example (Appel)

j
g
c
h




Graph Coloring Example (Appel)

f
k
j
g
c
h



Graph Coloring Example (Appel)

e
d
f
k
j
g
c
h



Graph Coloring Example (Appel)


b
m
e
d
f
k
j
g
c
h



Graph Coloring Example (Appel)


j
g
c
h

Well, gotta spill
(we might have
gotten lucky and still
found a color)




Graph Coloring Example (Appel)

g
c
h




Graph Coloring Example (Appel)




We need to load j from memory... into what?

```
r2 = mem[j + 12]
r1 = r1 - 1
r2 = r2 * r1
r3 = mem[j + 8]
r1 = mem[j + 16]
r2 = mem[r2]
r3 = r3 + 8
r3 = r3
r1 = r1 + 4
j = r2
```




Option 1: Move to a temp, do reg alloc again

```
temp1 = stack[0]
r2 = mem[temp1 + 12]
r1 = r1 - 1
r2 = r2 * r1
temp1 = stack[0]
r3 = mem[temp1 + 8]
temp1 = stack[0]
r1 = mem[temp1 + 16]
r2 = mem[r2]
r3 = r3 + 8
r3 = r3
r1 = r1 + 4
temp1 = r2
stack[0] = temp1
```



Option 2: Reserve a register or two for this


```
r4 = stack[0]
r2 = mem[r4 + 12]
r1 = r1 - 1
r2 = r2 * r1
r4 = stack[0]
r3 = mem[r4 + 8]
r4 = stack[0]
r1 = mem[r4 + 16]
r2 = mem[r2]
r3 = r3 + 8
r3 = r3
r1 = r1 + 4
r4 = r2
stack[0] = r4
```



Graph Coloring Example (Appel)

```
r4 = mem[r3 + 12]
r2 = r1 - 1
r2 = r4 * r2
r4 = mem[r3 + 8]
r1 = mem[r3 + 16]
r2 = mem[r2]
r3 = e + 8
r4 = r3 ←
r1 = r1 + 4
r3 = r2 ←
```


This
Next time:
Avoid these



Coalescing: Combining nodes to eliminate moves

```
g = mem[j + 12]
h = k - 1
f = g * h
e = mem[j + 8]
m = mem[j + 16]
b = mem[f]
c = e + 8
d = c
k = m + 4
j = b
```

Blue edge + no black edge: would like to coalesce




Coalescing unsafely can make a graph uncolorable

```
g = mem[j + 12]
h = k - 1
f = g * h
e = mem[j + 8]
m = mem[j + 16]
b = mem[f]
c = e + 8
d = c
k = m + 4
j = b
```


jb has all the edges from j and b!

We'd rather move than spill




Conservative coalescing strategies will always keep a graph colorable

- Briggs: a and b can be coalesced if the resulting node ab will have fewer than K neighbors of degree $\geq K$
 - (Recall: $K = \text{number registers/colors}$)




Conservative coalescing strategies will always keep a graph colorable

- Briggs: a and b can be coalesced if the resulting node ab will have fewer than K neighbors of degree $\geq K$
 - (Recall: $K = \text{number registers/colors}$)




Conservative coalescing strategies will always keep a graph colorable

- Briggs: a and b can be coalesced if the resulting node ab will have fewer than K neighbors of degree $\geq K$
 - (Recall: $K = \text{number registers/colors}$)




Conservative coalescing strategies will always keep a graph colorable

- Briggs is *conservative*:
 - Coalescing nodes following Briggs is guaranteed not to make a graph uncolorable
 - Briggs might miss nodes that could still be safely coalesced



Conservative coalescing strategies will always keep a graph colorable


- Briggs is *conservative*:
 - Coalescing nodes following Briggs is guaranteed not to make a graph uncolorable
 - Briggs might miss nodes that could still be safely coalesced



Conservative coalescing strategies will always keep a graph colorable

- George: Nodes a and b can be coalesced if, for every neighbor t of a , either:
 - t already interferes with b or
 - t has degree $< K$

j and b can be coalesced for $K=4$, not $K=3$




Conservative coalescing strategies will always keep a graph colorable

- George: Nodes a and b can be coalesced if, for every neighbor t of a , either:
 - t already interferes with b or
 - t has degree $< K$

j and b can be coalesced for $K=4$, not $K=3$


(and the graph is *not* 3-colorable!)




Graph coloring with coalescing

1. **Build** interference graph and classify nodes as move-related or non-move-related
2. **Simplify**, only removing non-move-related nodes of degree $< K$
3. **Coalesce** move-related nodes using a conservative heuristic
4. **Freeze** move-related nodes (give up trying to coalesce them) if can't simplify or coalesce
5. **Spill** (potentially) a node w/ degree $\geq K$, removing it from the graph and pushing it on the stack
6. **Select** colors for nodes in stack order


w/o coalescing:



w/ coalescing:




Coalescing Example (Appel)




Coalescing Example (Appel)

k
h
g




Coalescing Example (Appel)

k
h
g




Coalescing Example (Appel)

k
h
g




Coalescing Example (Appel)

cd
k
h
g



Coalescing Example (Appel)

jb
cd
k
h
g




Coalescing Example (Appel)

e
m
f
jb
cd
k
h
g




Coalescing Example (Appel)

jb
cd
k
h
g




Coalescing Example (Appel)

k
h
g




Coalescing Example (Appel)




Coalescing Example (Appel)

```
g = mem[j + 12]
h = k - 1
f = g * h
e = mem[j + 8]
m = mem[j + 16]
b = mem[f]
c = e + 8
d = c
k = m + 4
j = b
```




Coalescing Example (Appel)

```
r4 = mem[r1 + 12]
r2 = r2 - 1
r3 = r4 * r2
r4 = mem[r1 + 8]
r2 = mem[r1 + 16]
r1 = mem[r3]
r4 = r4 + 8
r4 = r4
r2 = m + 4
r1 = r1
```



Coalescing Example (Appel)

```
r4 = mem[r1 + 12]
r2 = r2 - 1
r3 = r4 * r2
r4 = mem[r1 + 8]
r2 = mem[r1 + 16]
r1 = mem[r3]
r4 = r4 + 8
r2 = m + 4
```



Another example

