


CS443: Compiler Construction

Lecture 21: Risc-V ISA

Stefan Muller

Based on material by Yan Garcia and Rujia Wang

You are here



An ISA is the set of instructions a computer can execute

- The job of a CPU
 - Fetch an instruction from memory
 - Decode
 - Execute
 - Write results to memory
 - Repeat (basically) forever

add x3, x2, x0


Assembler

01110001110110

There are many different ISAs with rich histories




By Mike Deerkoski - <https://www.flickr.com/photos/deerkoski/7178643521/in/photostream>



There are many different ISAs with rich histories

Intel

x86



x86

AMD

1990

x86-64

AMD64

2000

...

RISC (Reduced Instruction Set Computer)

idea: simpler, faster hardware

- Earlier philosophy (“CISC”):
Want to do something new? Add an instruction!
- RISC: Cocke, Hennessy, Patterson (1980s)

RISC-V: A simple RISC Architecture, good for teaching




- Originally developed in 2010 at UC Berkeley for teaching
- Open-source

Assembly Language: Human-readable machine code

- Assembly language is tied to ISA
- (Roughly) 1-to-1 correspondence with ISA instructions
 - (Some assembly languages offer convenient mnemonics that expand to multiple instructions)

An instruction is an opcode and operands
(registers)



- Operands can only be **registers** and sometimes constants (“immediates”)
- Registers: Limited number of single-word storage locations in hardware

Registers in RISC-V

- (Also some floating point registers we won't talk about)

Register	ABI Name
x0	zero
x1	ra
x2	sp
x3	gp
x4	tp
x5-7	t0-2
x8	s0/fp
x9	s1
x10-11	a0-1
x12-17	a2-7
x18-27	s2-11
x28-31	t3-t6

Before we dive into RISC-V: A quick recap on data representation

- Bit (binary digit): 0 or 1
- “Nibble”: 4 bits (1 hex digit 0x0-0xF)
- Byte: 8 bits
 - 2 hex digits: 0x00-0xFF
- Word: “Natural” size of data operated on by a computer
 - 32-bit ISA: 32 bits (4 bytes)
 - Width of registers

Integers in binary/hex

1	0	1	0			
$\times 2^3$	+	$\times 2^2$	+	$\times 2^1$	+	$\times 2^0$
2		2		10		
10		a		1010		
16		10		10000		
32		20		100000		

← →

“Most significant” “Least significant”

Review: Endianness

- Store data one byte at a time
 - Order of bits in a byte doesn't change!
- So do we store the most significant byte at the lowest memory address (the way we'd write it left-to-right) or the highest?
 - Lowest: “Big-endian” (e.g., IBM System/360)
 - Highest: “Little-endian” (e.g., x86, RISC-V)

Little-endian

- 0xdeadbeef

ef	be	ad	de
----	----	----	----

Two's complement signed integers

- A 1 in MSB (Most significant bit) subtracts 2^{31} (instead of adding it)
- $100000\dots = -2^{31}$
- $011111\dots = 2^{31}-1$ (highest positive # representable)
- $111111\dots = -1$
- Can just add two's complement #'s without casing on sign!

Two's complement means two ways to extend integers to the left

1010101
←

- If signed int: want to sign-extend (extend with MSB)
 - LLVM: sext
 - 101 as 3-bit int = -3 = 11101 as 5-bit int
- If unsigned: want to zero-extend (extend with 0s)

Assembly operands, registers are untyped

- Value is whatever we interpret it as – (signed/unsigned) int/char/bool, etc.

add x3, x2, x1

Overflow:

char: Yes. unsigned int: No. signed int: Yes.

Still want types? Never fear




TALx86: A Realistic Typed Assembly Language*

Greg Morrisett Karl Crary[†] Neal Glew Dan Grossman Richard Samuels
Frederick Smith David Walker Stephanie Weirich Steve Zdancewic
Cornell University

1999



Registers are inside the processor



Q: Why not make a bigger processor with more registers?

RISC-V Instructions are 32 bits

- 6 types of instructions:

31	30	25 24	21	20	19	15 14	12 11	8	7	6	0	
funct7		rs2		rs1		funct3		rd		opcode		R-type
		imm[11:0]		rs1		funct3		rd		opcode		I-type
	imm[11:5]		rs2		rs1	funct3		imm[4:0]		opcode		S-type
imm[12]	imm[10:5]		rs2		rs1	funct3	imm[4:1]	imm[11]	opcode			B-type
		imm[31:12]					rd		opcode			U-type
imm[20]		imm[10:1]		imm[11]		imm[19:12]		rd		opcode		J-type

R-type instruction: Destination, two register operands

Risc-V	LLVM	C
add x1, x2, x3	%x1 = add i32 %x2 %x3	x1 = x2 + x3
sub x3, x4, x5	%x3 = sub i32 %x4 %x5	x3 = x4 - x5

Also: xor, or, and, mul, div
divu (div unsigned)
sll (shift left logical)
srl (shift right logical) – fill left with 0s
sra (shift right arithmetic) – fill left with sign bit
slt (set rd to 1 iff rs1 < rs2)

x_0 is always 0, writes are ignored

- Why would you want to read from x_0 ?
 - $mv\ rd, rs = add\ rd\ rs\ x_0$
- Why would you want to write to x_0 ?
 - $nop = add\ x_0\ x_0\ x_0$
 - (There are other ways to write a no-op instruction, but this is the conventional one)

I-type instructions: Destination, register, immediate

Risc-V

addi x1, x2, n

subi x3, x4, n

LLVM

%x1 = add i32 %x2 n

%x3 = sub i32 %x4 n

C

x1 = x2 + n

x3 = x4 - n

Also: xori, ori, andi, (NO muli, divi)

slti (set if less-than)

slli (shift left logical)

srl (shift right logical) – fill left with 0s

srai (shift right arithmetic) – fill left with sign bit

Example

```
%x = mul i32 %y 2
```

$x_1 \leftarrow x$ $x_2 \leftarrow y$

- add x_1, x_2, x_2
- slli $x_1, x_2, 1$
- addi $x_1, x_0, 2$
 $\text{mul } x_1, x_2, x_1$

Remember: You only get 12 bits for immediate (not very big)

- In RISC-V immediates are "sign extended"
 - So the upper bits are the same as the largest bit
 - Remember sign extended 2's complement..
- So for a 12b immediate...
 - Bits 31:12 get the same value as Bit 11

31	30	25 24	21	20	19	15 14	12 11	8	7	6	0	
		funct7		rs2		rs1	funct3		rd		opcode	R-type
		imm[11:0]			rs1	funct3		rd		opcode		I-type

If you need big immediates, need 2 insts

Risc-V

lui x1, n

C

x1 = n << 12 (x1 = n * 4096)

%x = add i32 %y, 5000

x1 <- x x2 <- y

5000 = 1 0011 1000 1000

lui x1, 1

addi x1, x1, 904

add x1, x1, x2

Control flow in RISC-V: similar to LLVM, but less structured

Assembly:

loopforever:

```
add x0, x0, x0  
j loopforever
```

After assembling/linking:

```
add x0, x0, x0  
j -4
```

Offset: Position independent

j isn't actually an instruction

- It's a “pseudoinstruction” that gets expanded into other instructions by the assembler (like mv, nop)

B-type instructions (Conditional branches): 2 registers and a label/offset

Risc-V

beq x1, x2, addr

LLVM

```
%x3 = icmp eq i32 %x1 %x2  
br i1 %x3, label addr, ???
```

C

```
if (x1 == x2) goto addr
```

Also: bne, blt, bge, (bltu, bgeu)

NO ble, bgt

Example

```
%x1 = icmp lt i32 %x2, %x3  
br i1 %x1, label ltrue, label lfalse
```

```
blt x2, x3, ltrue  
j lfalse
```

```
slt x1, x2, x3  
bne x1, x0, ltrue  
j lfalse
```

Unlike LLVM, control “falls through” to next instruction

Example

```
%x1 = icmp le i32 %x2, %x3  
br i1 %x1, label ltrue, label lfalse  
  
bge x3, x2, ltrue  
j lfalse
```

Example

Assuming assignments below, compile *if* block

$f \rightarrow x_{10}$ $g \rightarrow x_{11}$ $h \rightarrow x_{12}$


$i \rightarrow x_{13}$ $j \rightarrow x_{14}$

if (i == j)	bne x13,x14,done
f = g + h;	add x10,x11,x12
done:	

Unconditional jump instructions: jal, jalr

- **jal rd, imm**
 - Jump to label (or by offset)
 - Set $rd = PC + 4$ (next instruction after jal)
- **jalr rd, rs, imm**
 - Jump to address in $rs + imm$
 - Set $rd = PC + 4$ (next instruction after jal)
- $j imm = jal x0, imm$

Loading from and storing to memory



3
1

Memory is addressed in bytes

- (But access memory a word at a time, so in practice, will only access memory at multiples of 4 bytes)
- Generally: data ≥ 1 word must be *aligned* to addresses that are multiples of 4

lw loads from memory to register

`lw rd, imm(rs)`

Load word at $rs + imm$ into rd

lw loads from memory to register

C code

```
int A[100];  
g = h + A[3];
```

Register, register,
immediate: lw is an
I-type instruction

Using Load Word (lw) in RISC-V:

- `lw x10,12(x13) # Reg x10 gets A[3]`
- `add x11,x12,x10 # g = h + A[3]`
- Assume: x13 – base register (pointer to A[0]) Note: 12 – offset in bytes
- Offset must be a constant known at *assembly time*

sw transfers from register to memory

C

```
int A[100]  
A[10] = h + A[3]
```

RISC-V

```
lw x10, 12(x13)  
add x10, x12, x10  
sw x10, 40(x13)
```

Note:

- x13 – base register (pointer)
- 12, 40 – offsets in bytes
- x13 + 12 and x13 + 40 must be multiples of 4 to maintain alignment

Example

```
addi x11,x0,0xfeed  
addi x12,x0,0xbeef  
addi x6,x5,4  
sw x11,0(x5)  
sw x12,4(x5)  
lw x12,0(x6)
```

- What's the value in x12? Answer: 0xbeef


Example

```
addi x11,x0,0xfeed  
addi x12,x0,0xbeef  
addi x6,x5,1  
sw x11,0(x5)  
sw x12,4(x5)  
lw x12,0(x6)
```

- What's the value in x12?

Answer: Undefined

Memory layout in RISC-V



A stack frame is where we store spilled locals, plus anything `alloca`'d

