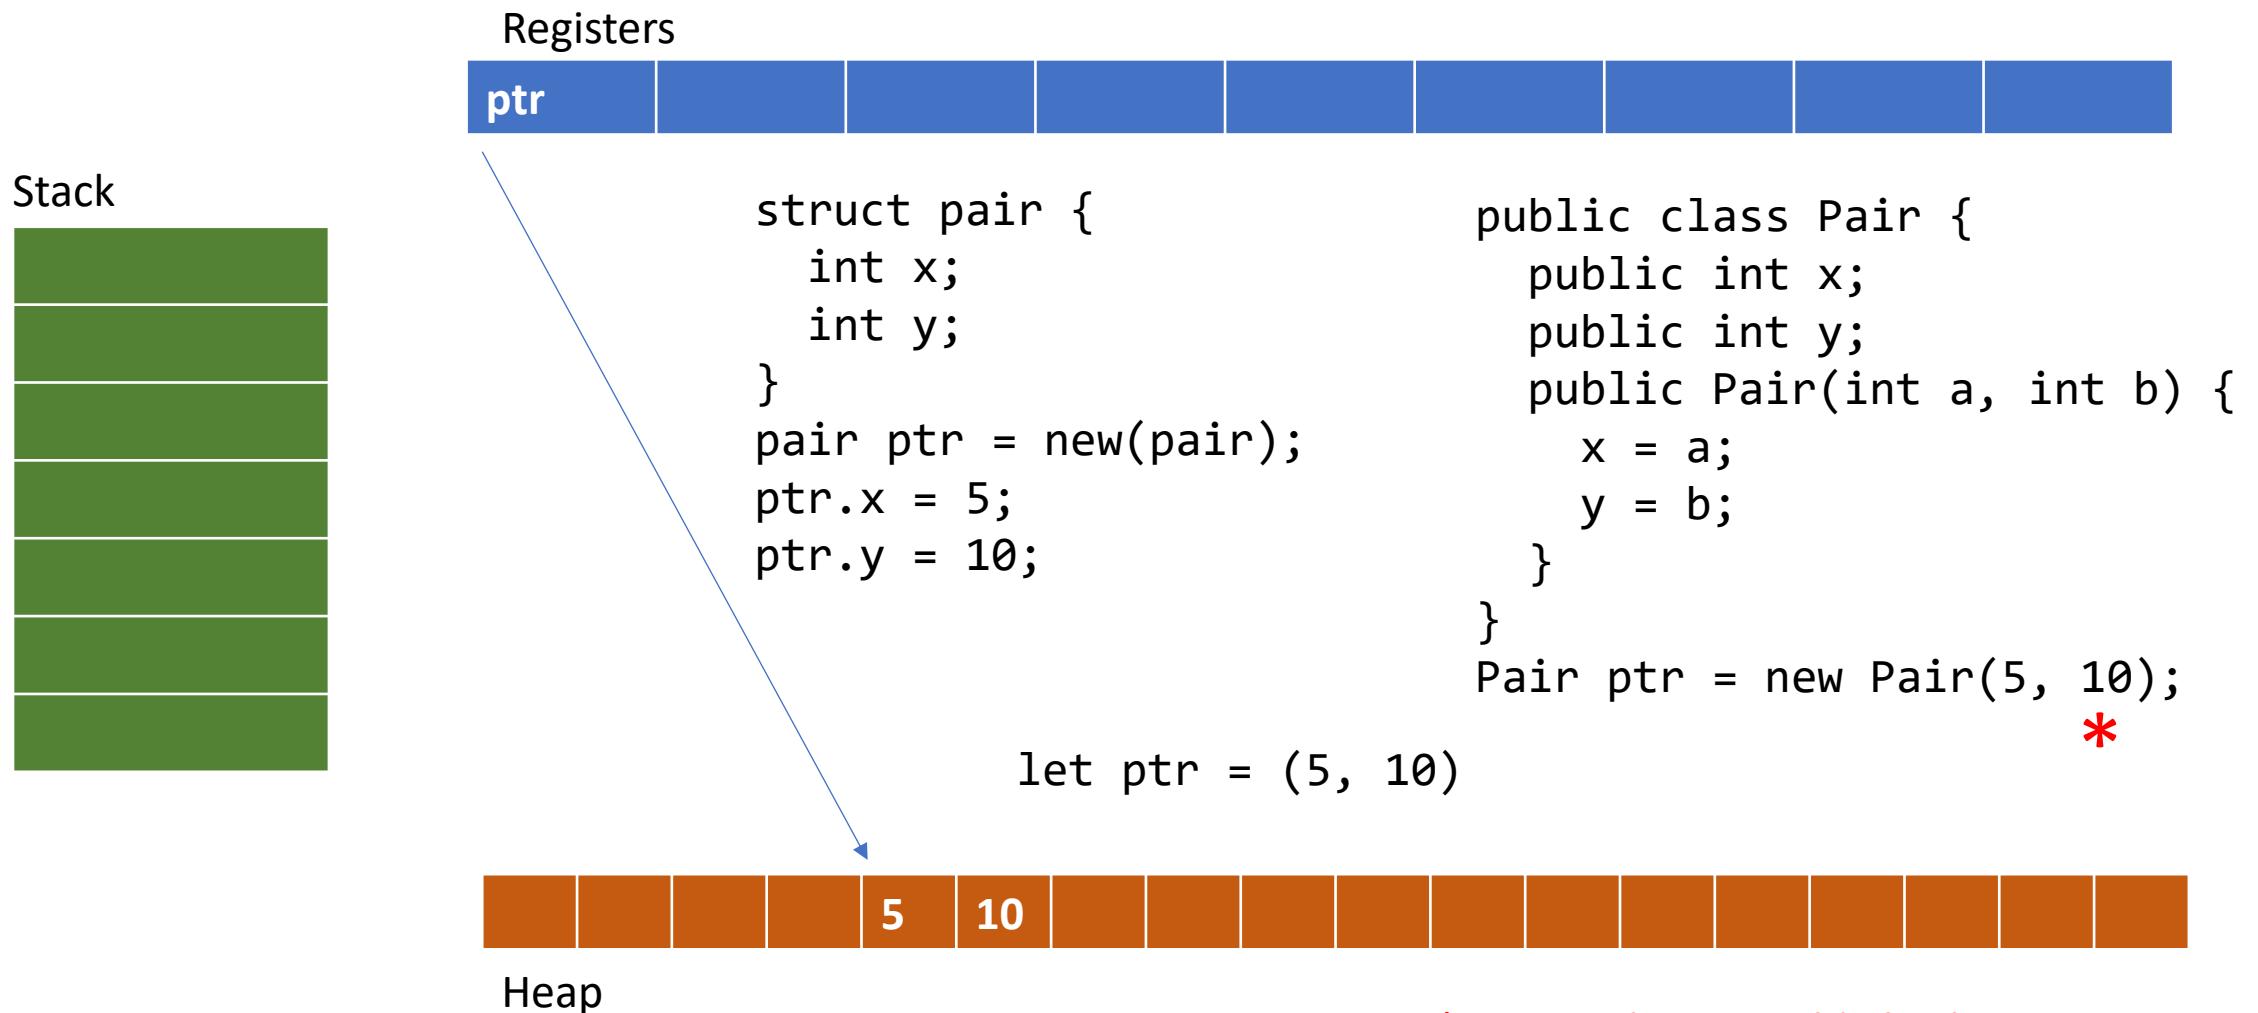


CS443: Compiler Construction

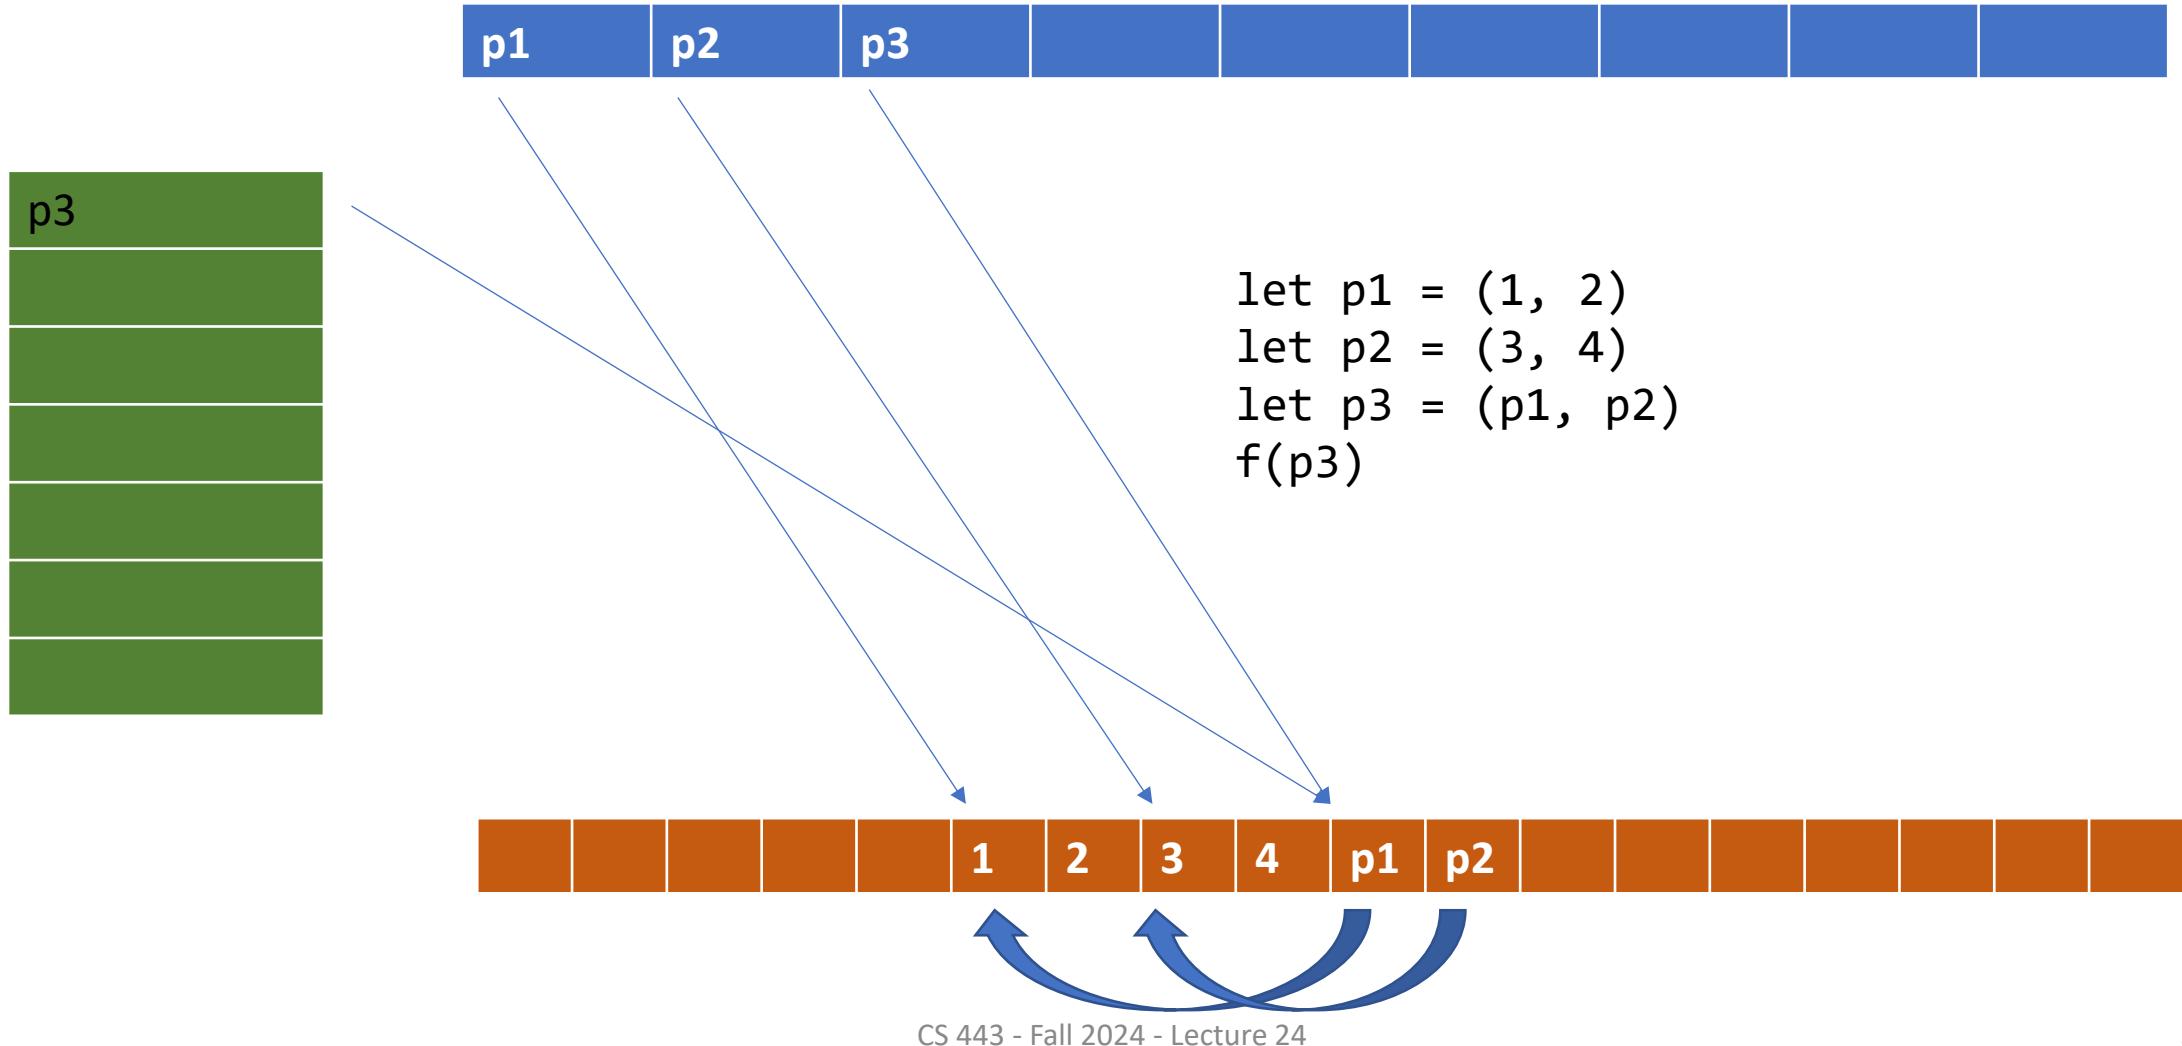
Lecture 24: Memory Management & Garbage Collection

Stefan Muller

Memory layout



Objects can be nested



Memory management answers two questions

- How do we allocate memory?
- What do we do with it when we're done?

MM breaks down into two basic strategies

- Manual – programmer says when to allocate (`malloc/new`) and free (`free/drop`)
 - Good control
 - Might forget to free/free twice/use after free
- Automatic – free memory automatically when no longer needed
 - (“Garbage collection”)
 - Some runtime overhead

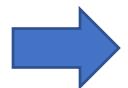
What about Rust?

- Still manual, the compiler just inserts calls to drop when variables go out of scope (definitely can't be used any more)
- Overly conservative, but prevents errors with free.
- Manual doesn't have to mean awful!

Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```



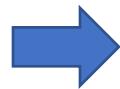
```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```



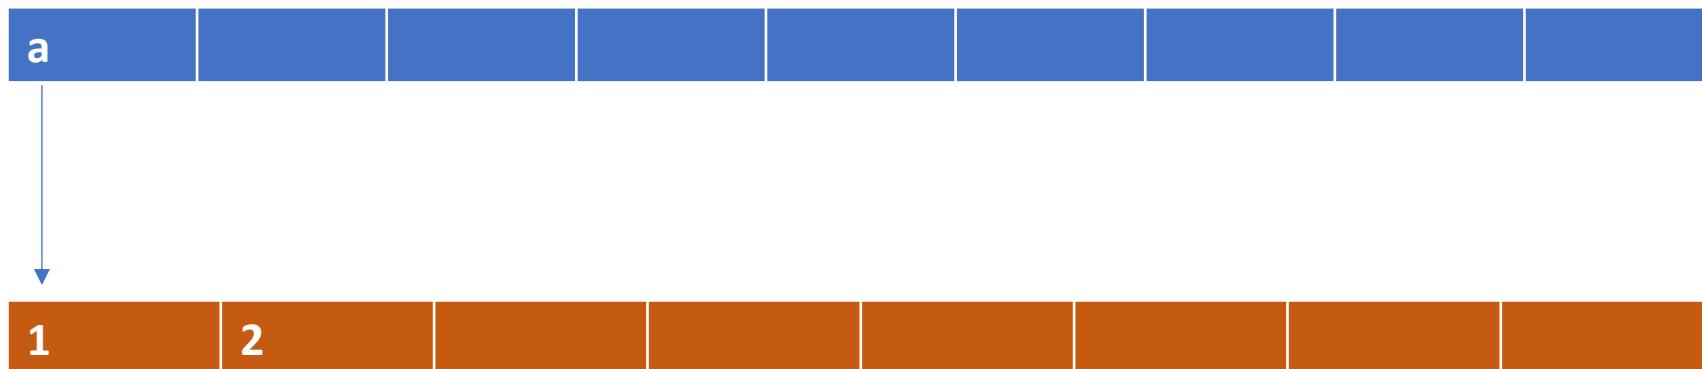
Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```



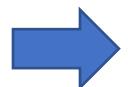
```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```



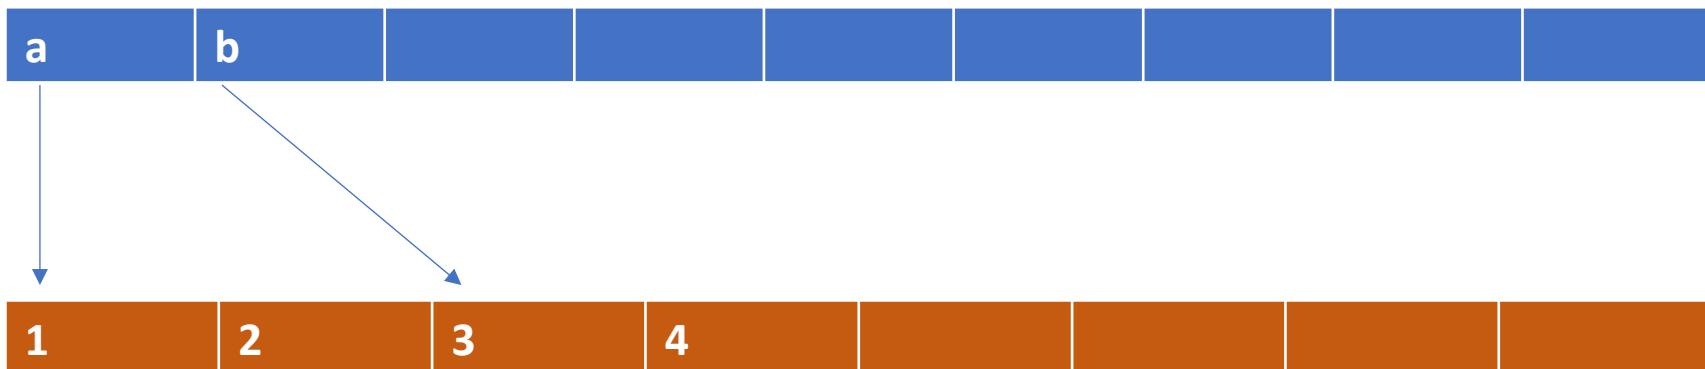
Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```



```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```

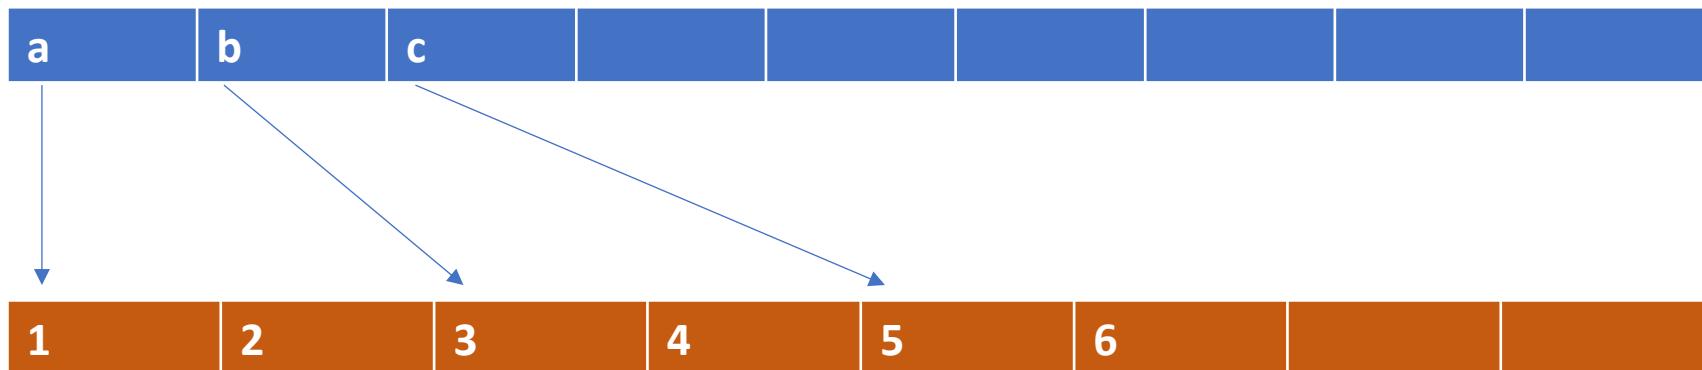


Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```

```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```

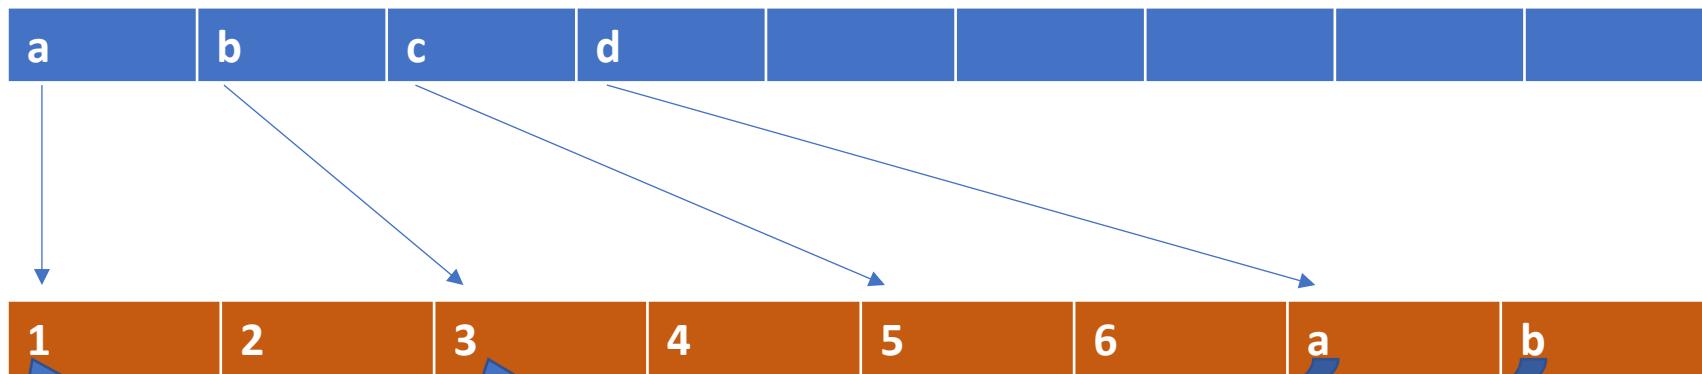
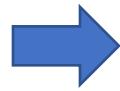


Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```

```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```

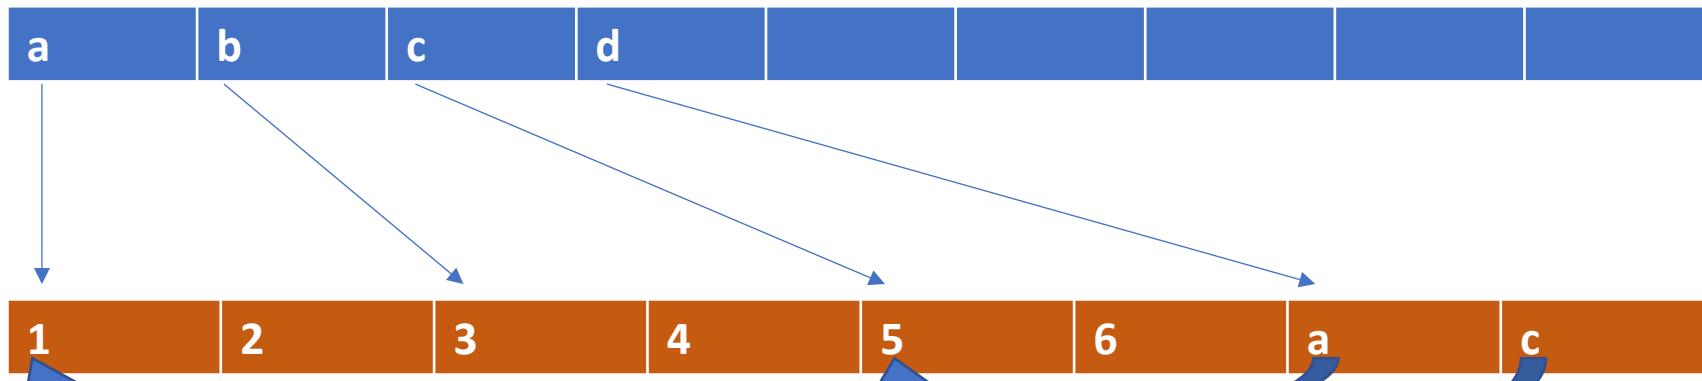


Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```

```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```



Manual Memory Management

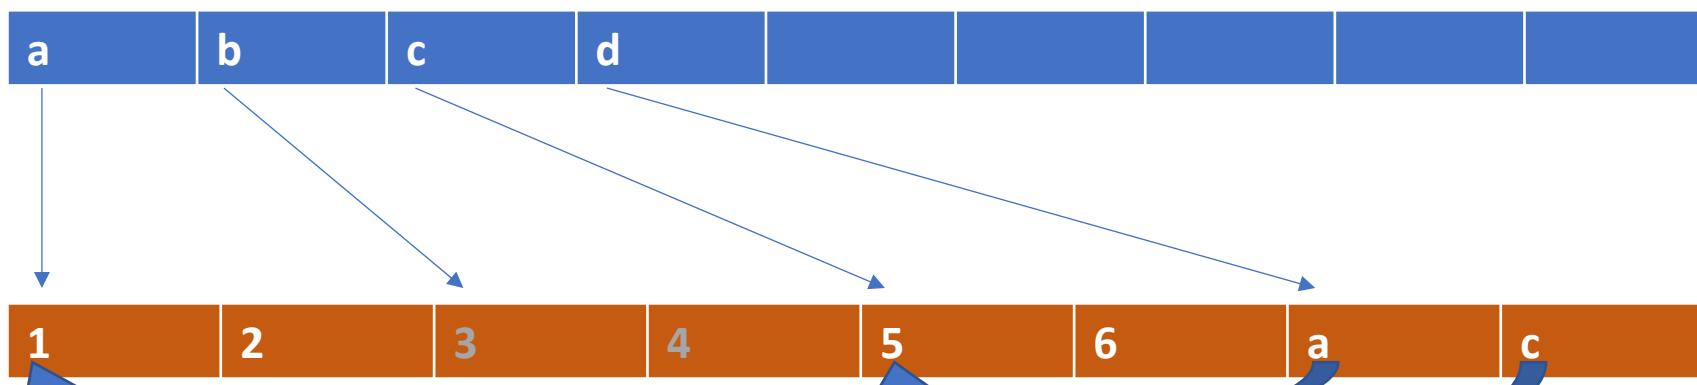
Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```

```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```



b still valid,
still points
to same loc,
but can
reuse
memory



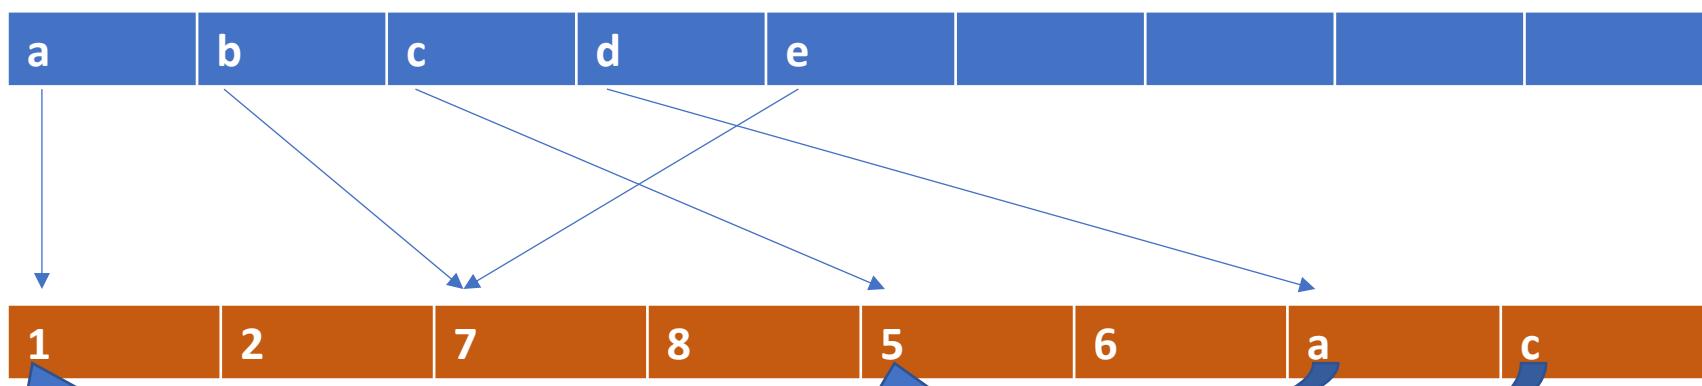
Manual Memory Management

Let

```
pair a = pair(x, y);  
mean  
pair a =  
malloc(sizeof(pair));  
a.fst = x;  
a.snd = y;
```

Need to reuse
that space—
fragmentation

```
pair a = pair(1, 2);  
pair b = pair(3, 4);  
pair c = pair(5, 6);  
pair d = pair(a, b);  
d.snd = c;  
free(b);  
pair e = pair(7, 8);
```



Manual pros and cons

- Pros
 - Space-efficient
 - `free` is cheap
 - Lots of control
- Cons
 - `malloc` is expensive (and hard to implement)!
 - Lots of control

Functional languages allocate *a lot*

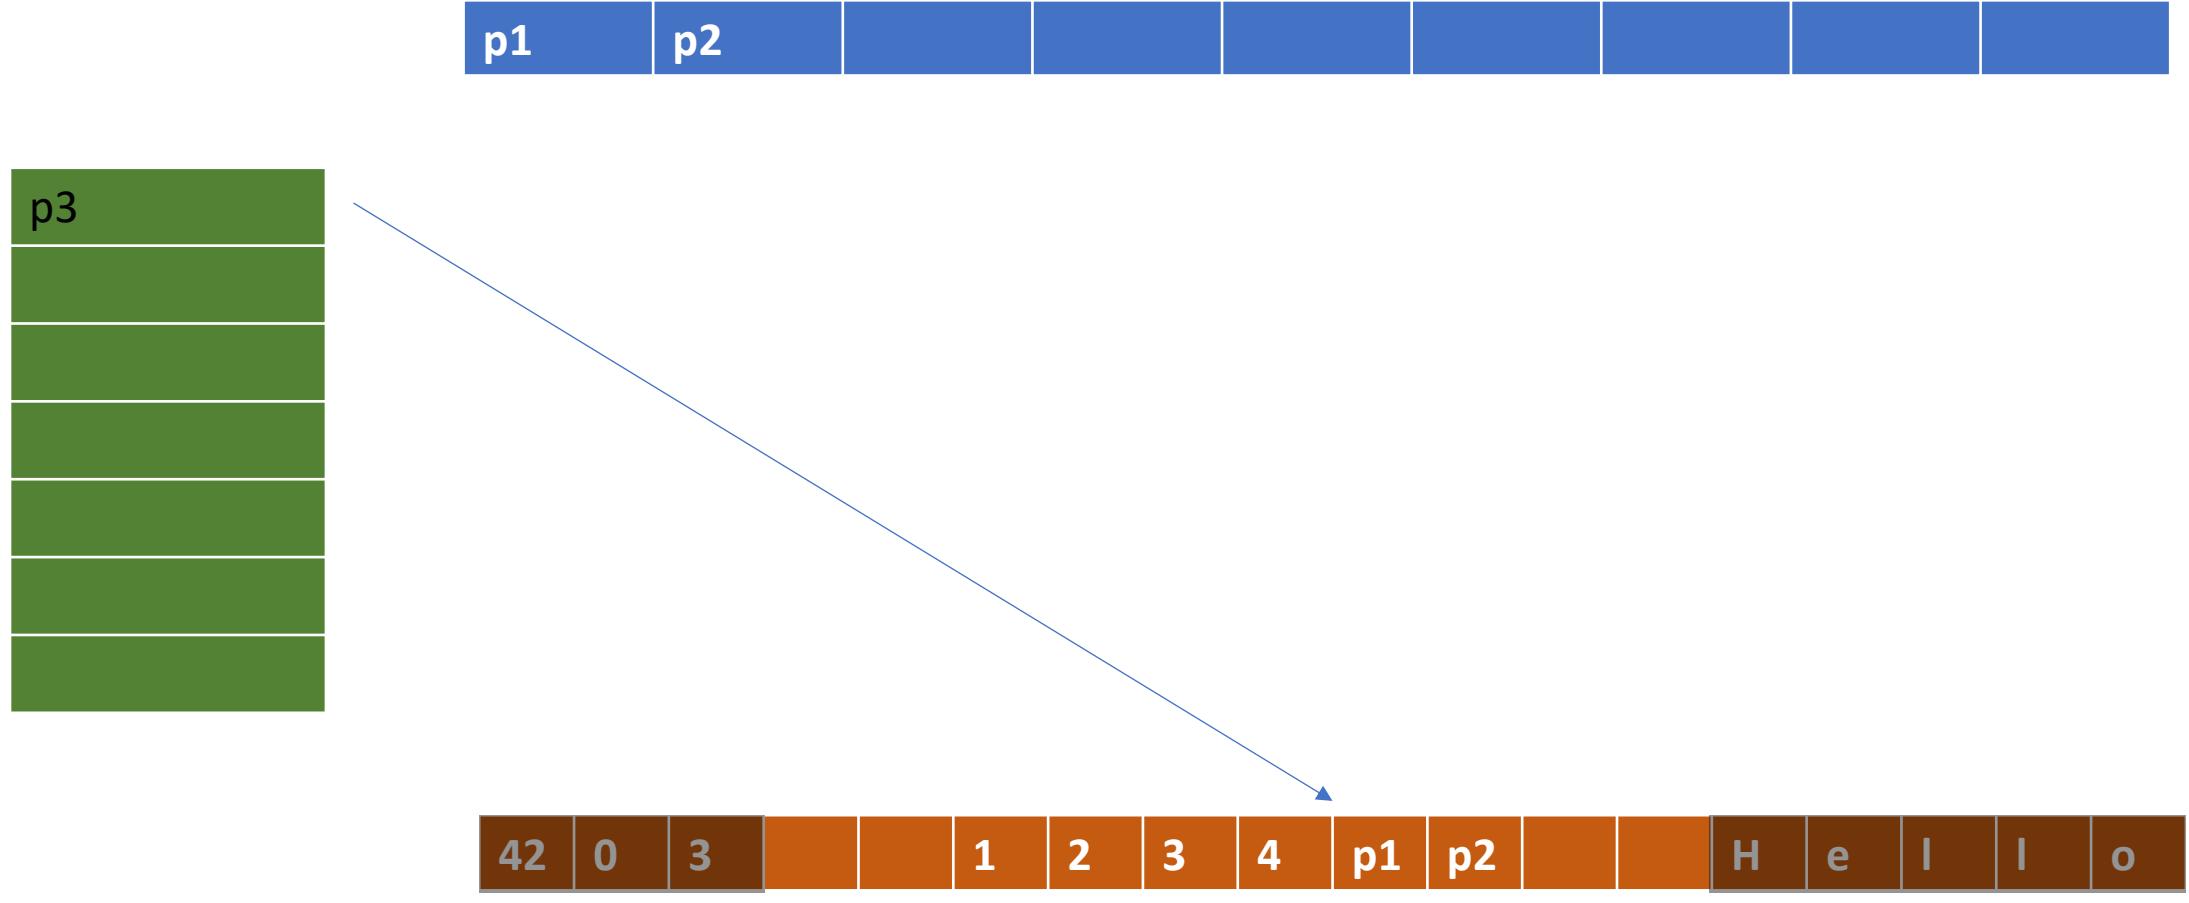
```
__list __ctemp52 = new(__list);
__ctemp52.list_tl = __env;
__ctemp52.list_hd = ((int)(__list) __ctemp51.list_hd);
__env = __ctemp52;
__list __ctemp53 = new(__list);
__ctemp53.list_tl = __env;
__ctemp53.list_hd = ((int) __ctemp51.list_tl);
__env = __ctemp53;
__clos __ctemp54 = ((__clos) __lookup(4, __env));
__clos __ctemp55 =
  ((__clos(*) (int, __list)) __ctemp54.clos_fun) (((int) __lookup(1, __env)),
  __ctemp54.clos_env);
__clos __ctemp56 = __ctemp55;
__pair __ctemp57 =
  ((__pair(*) (__list, __list)) __ctemp56.clos_fun) (((__list) __lookup(0,
  __env)), __ctemp56.clos_env);
__pair __ctemp58 = __ctemp57;
__list __ctemp59 = new(__list);
__ctemp59.list_tl = __env;
__ctemp59.list_hd = ((int) __ctemp58.pair_fst);
__env = __ctemp59;
__list __ctemp60 = new(__list);
__ctemp60.list_tl = __env;
__ctemp60.list_hd = ((int) __ctemp58.pair_snd);
```

And also can you imagine having to free everything manually?

Reachability and garbage

- Root set: Anything immediately reachable (registers, stack)
 - e.g., local variables, arguments
- Reachable (“live”): any objects (transitively) pointed to by root set
- Garbage (“dead”): any allocated objects not reachable

Objects not reachable from roots are dead/garbage



Knowing what points to what isn't as easy as it sounds

- In C:

```
int *p = (int *)0xdeadbeef;  
*p = 5;
```

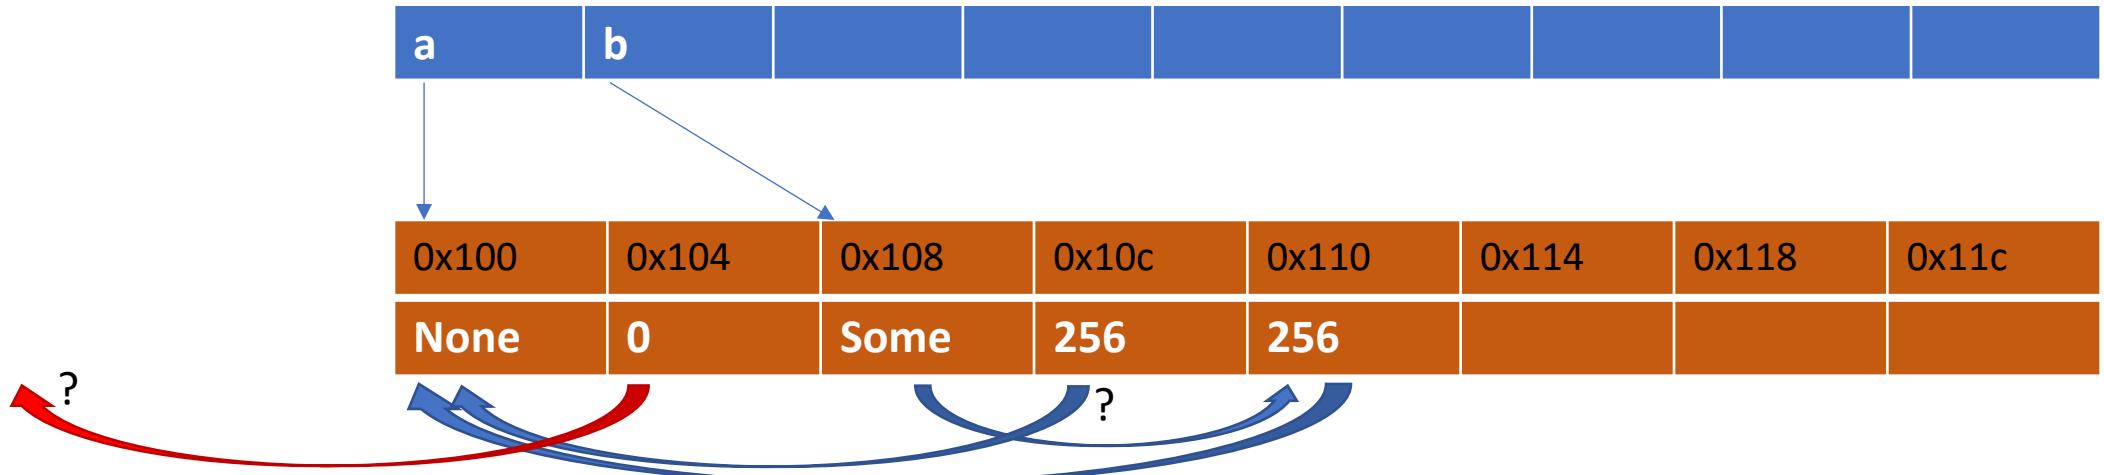


Garbage collection won't work well in C

Knowing what points to what isn't as easy as it sounds

In ML

```
let a = (None, 0)  
let b = (Some a, 256)
```



OCaml's clever hack: use the LSB to indicate integer or pointer



LSB of a ptr is 0 anyway

(1, 2, x + 1)

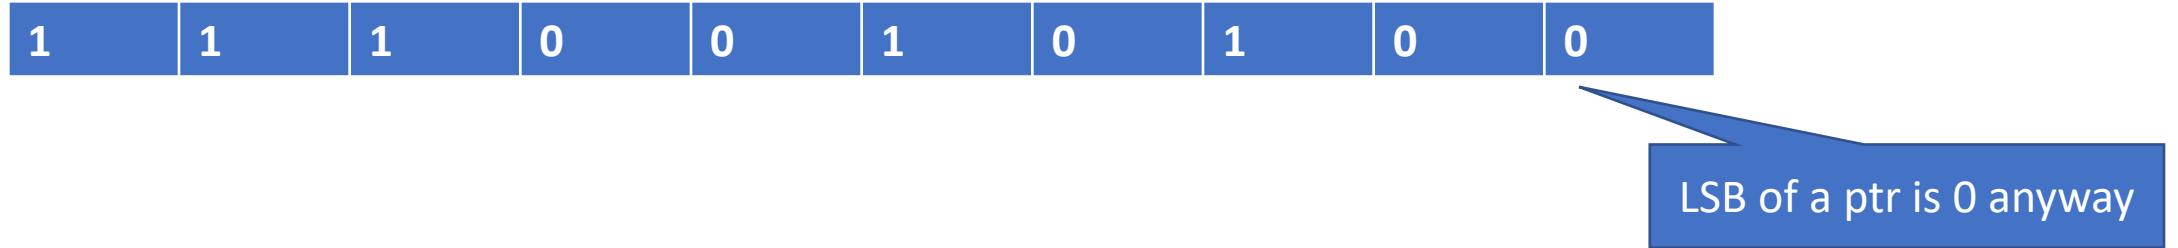
OCaml's clever hack: use the LSB to indicate integer or pointer



LSB of a ptr is 0 anyway

(1, 2, x + 1)	12d43:	48 c7 00 03 00 00 00	movq	\$0x3,(%rax)
	12d4a:	48 c7 40 08 05 00 00	movq	\$0x5,0x8(%rax)
	12d51:	00		
	12d52:	48 83 c3 02	add	\$0x2,%rbx
	12d56:	48 89 58 10	mov	%rbx,0x10(%rax)
	12d5a:	48 83 c4 08	add	\$0x8,%rsp
	12d5e:	c3	retq	

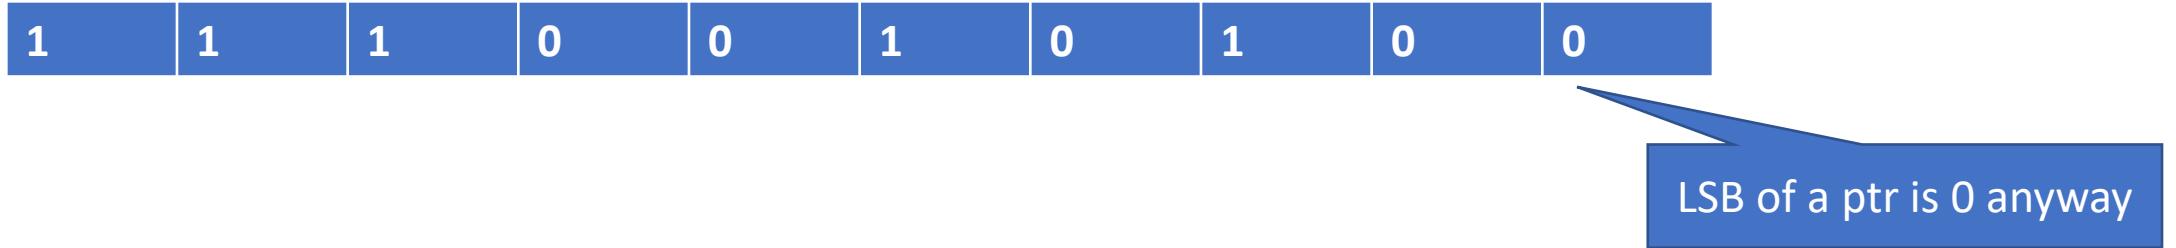
OCaml's clever hack: use the LSB to indicate integer or pointer



$$3 = 1 \ll 1 + 1$$

(1, 2, x + 1)	12d43:	48 c7 00 03 00 00 00	movq	\$0x3, (%rax)
	12d4a:	48 c7 40 08 05 00 00	movq	\$0x5, 0x8(%rax)
	12d51:	00		
	12d52:	48 83 c3 02	add	\$0x2,%rbx
	12d56:	48 89 58 10	mov	%rbx, 0x10(%rax)
	12d5a:	48 83 c4 08	add	\$0x8,%rsp
	12d5e:	c3	retq	

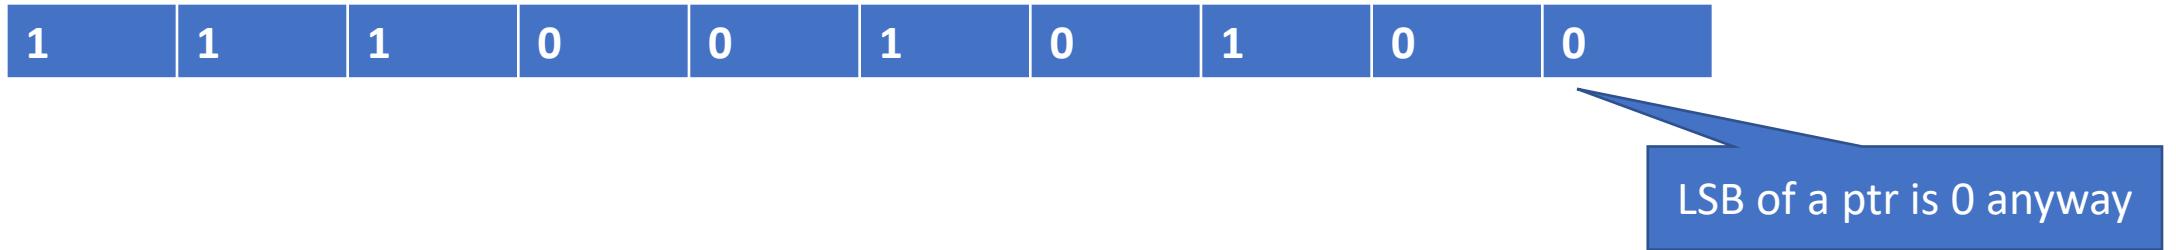
OCaml's clever hack: use the LSB to indicate integer or pointer



$$5 = 2 \ll 1 + 1$$

(1, 2, x + 1)	12d43:	48 c7 00 03 00 00 00	movq	\$0x3, (%rax)
	12d4a:	48 c7 40 08 05 00 00	movq	\$0x5, 0x8(%rax)
	12d51:	00		
	12d52:	48 83 c3 02	add	\$0x2,%rbx
	12d56:	48 89 58 10	mov	%rbx, 0x10(%rax)
	12d5a:	48 83 c4 08	add	\$0x8,%rsp
	12d5e:	c3	retq	

OCaml's clever hack: use the LSB to indicate integer or pointer

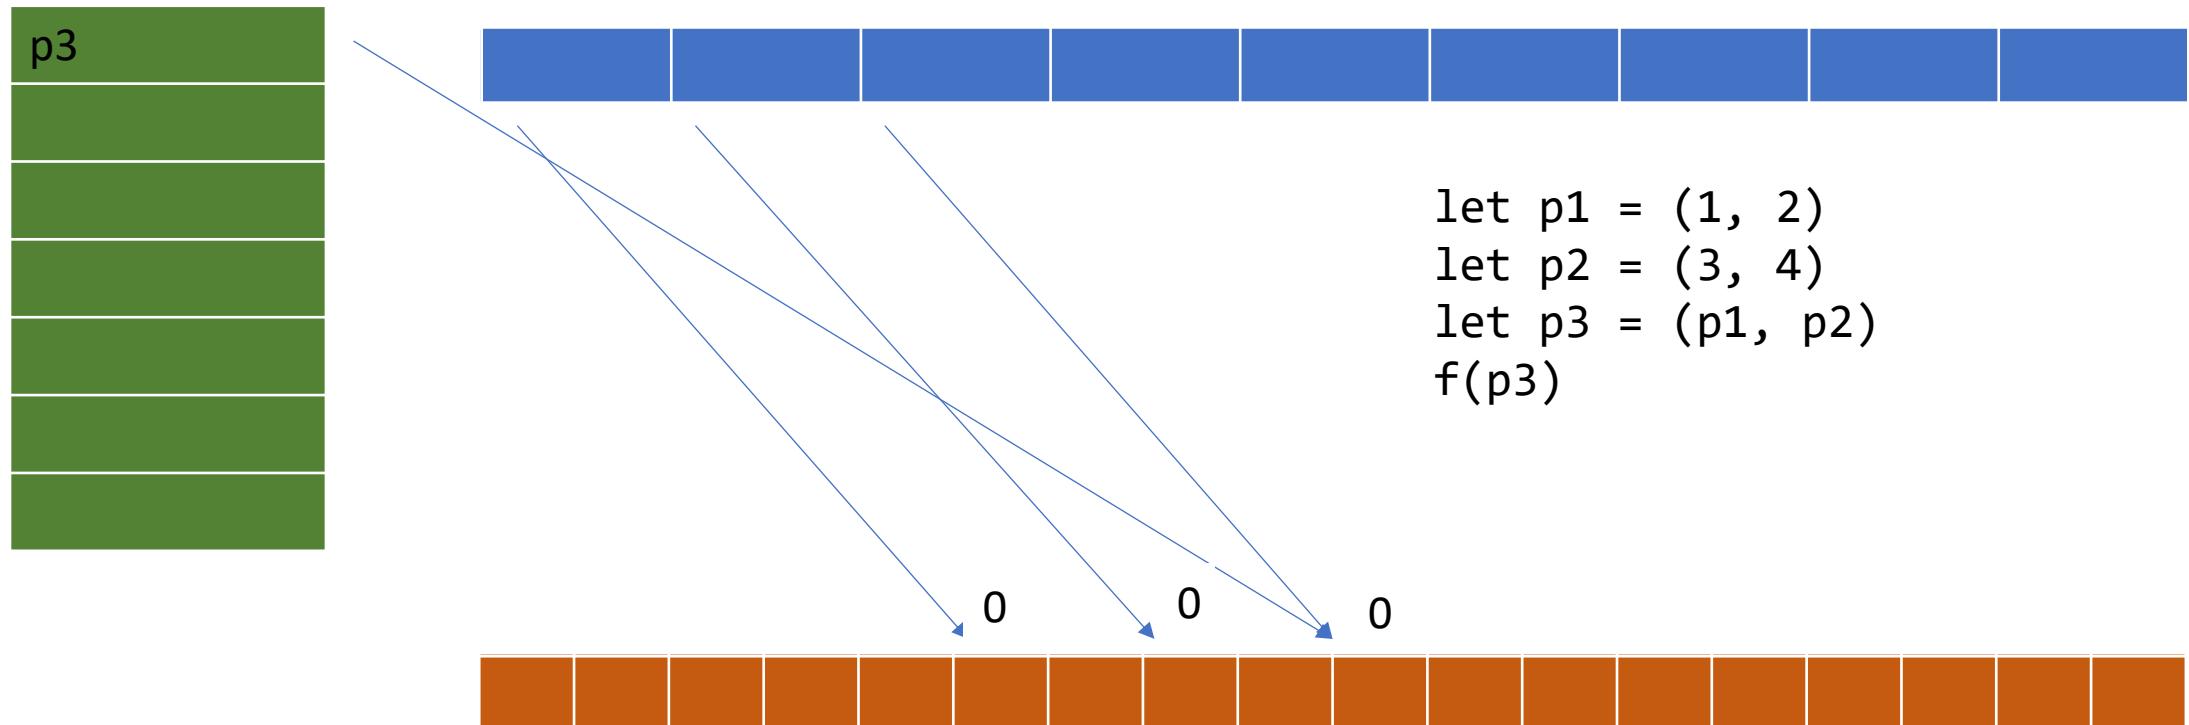


$$(x \ll 1 + 1) + (1 \ll 1) = (x + 1) \ll 1 + 1$$

(1, 2, x + 1)	12d43:	48 c7 00 03 00 00 00	movq	\$0x3, (%rax)
	12d4a:	48 c7 40 08 05 00 00	movq	\$0x5, 0x8(%rax)
	12d51:	00		
	12d52:	48 83 c3 02	add	\$0x2,%rbx
	12d56:	48 89 58 10	mov	%rbx, 0x10(%rax)
	12d5a:	48 83 c4 08	add	\$0x8,%rsp
	12d5e:	c3	retq	

GC Strategy #1: Reference counting

- Idea: keep track of how many references every object has

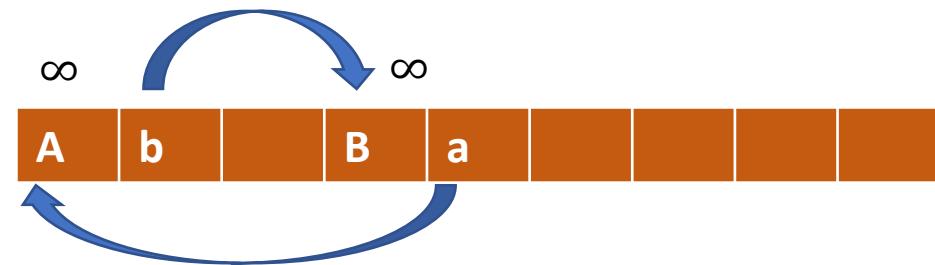


Reference counting pros

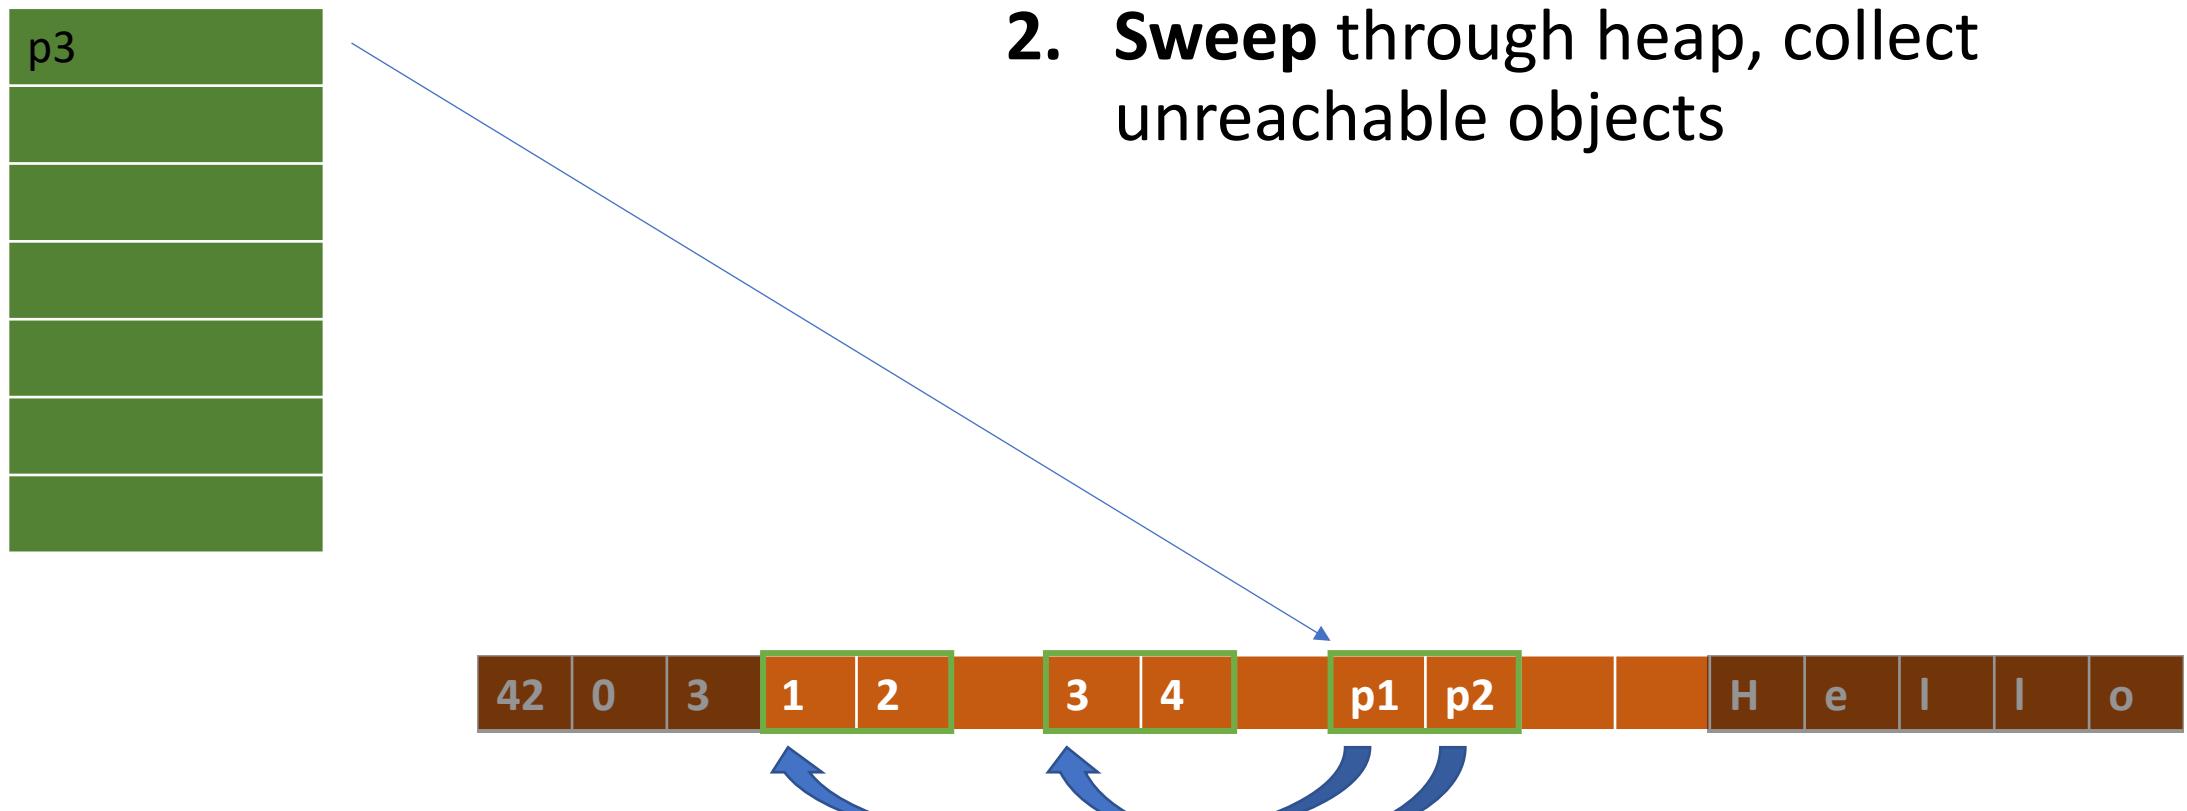
- Simple, intuitive
- Garbage collected immediately

Reference counting cons

- Cyclic data structures `a = new A();
b = new B();
A.b = b;
B.a = a;`
- Updating counts can be expensive



GC Strategy #2: Mark and sweep

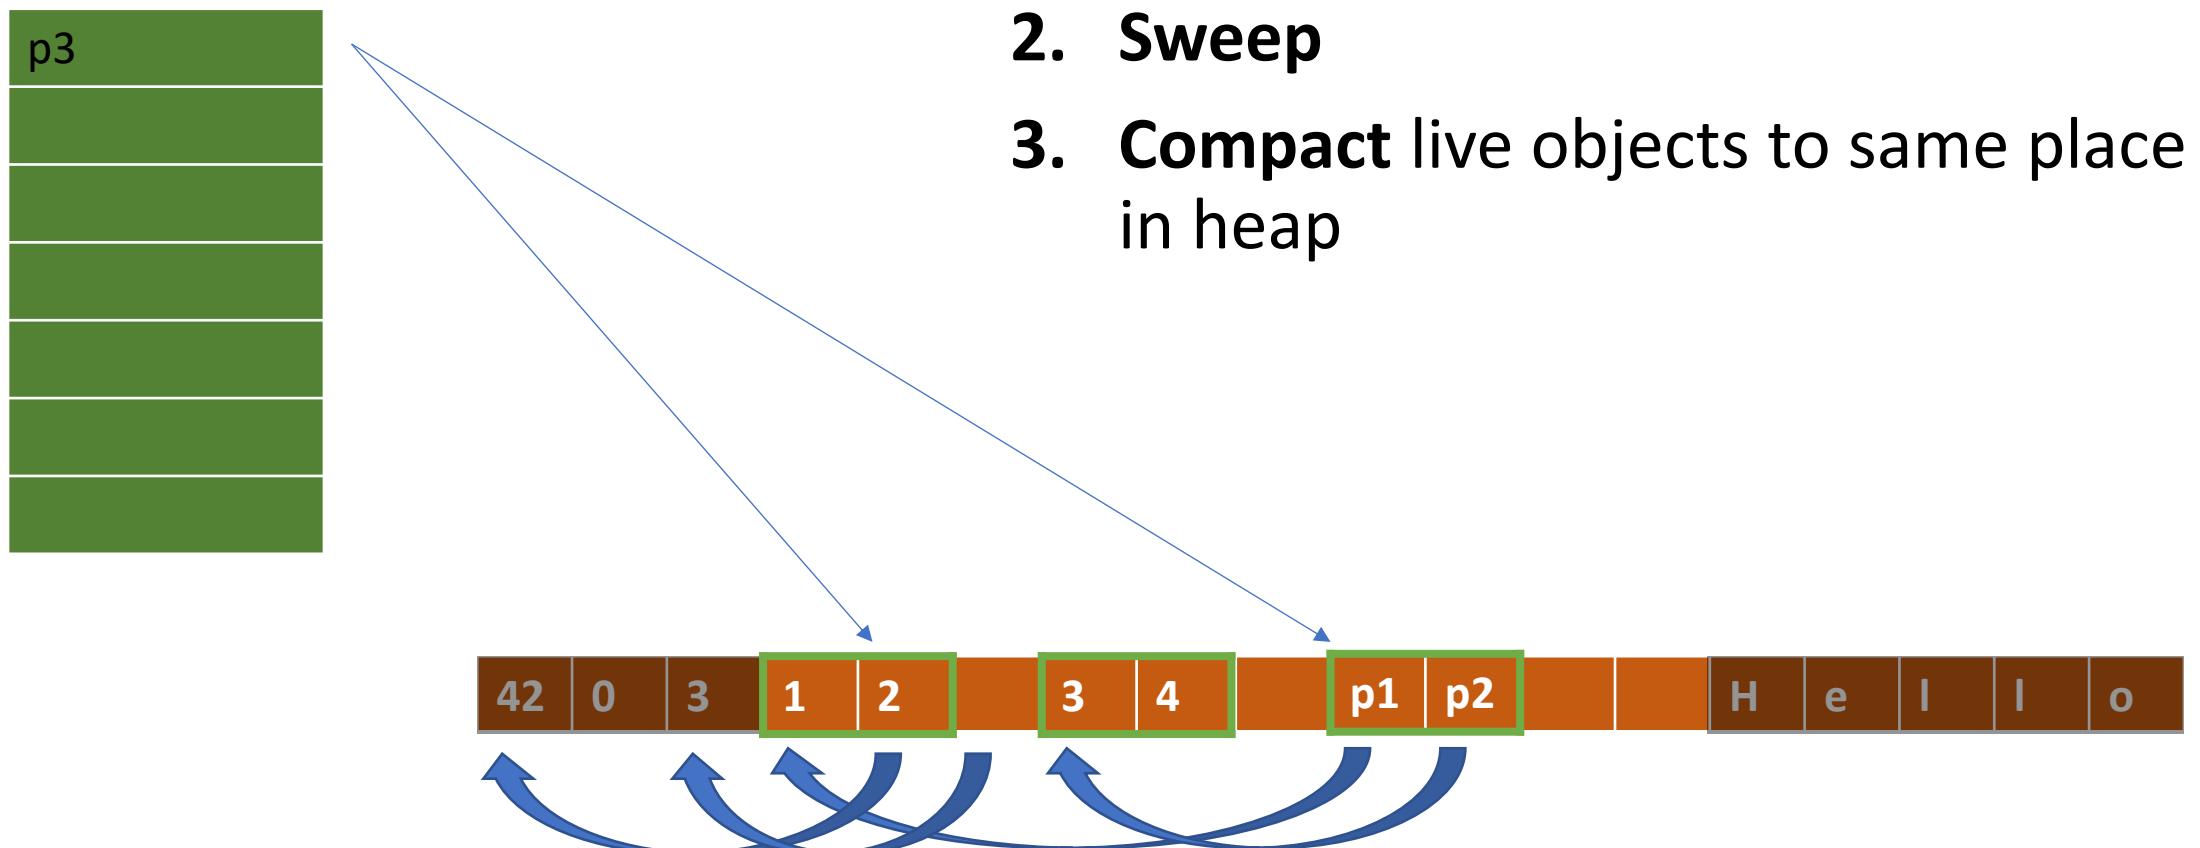


Mark and Sweep pros and cons

- Pros:
 - Works on cyclic references
 - Just traverse references once
- Cons:
 - Have to sweep through whole heap (can optimize)
 - Fragmentation



GC Strategy #2½: Mark and compact



Mark and compact pros and cons

- Pros:
 - Fragmentation solved
- Cons:
 - Have to update pointers

Implementing Compaction (#1): Keep a “forwarding pointer” in each object

1. Compute new locations of objects
2. Update all pointers
3. Move



Implementing Compaction (#1): Keep a “forwarding pointer” in each object

1. Compute new locations of objects
2. Update all pointers
3. Move



Implementing Compaction (#1): Keep a “forwarding pointer” in each object

1. Compute new locations of objects
2. Update all pointers
3. Move



Implementing Compaction (#1): Keep a “forwarding pointer” in each object

1. Compute new locations of objects
2. Update all pointers
3. Move



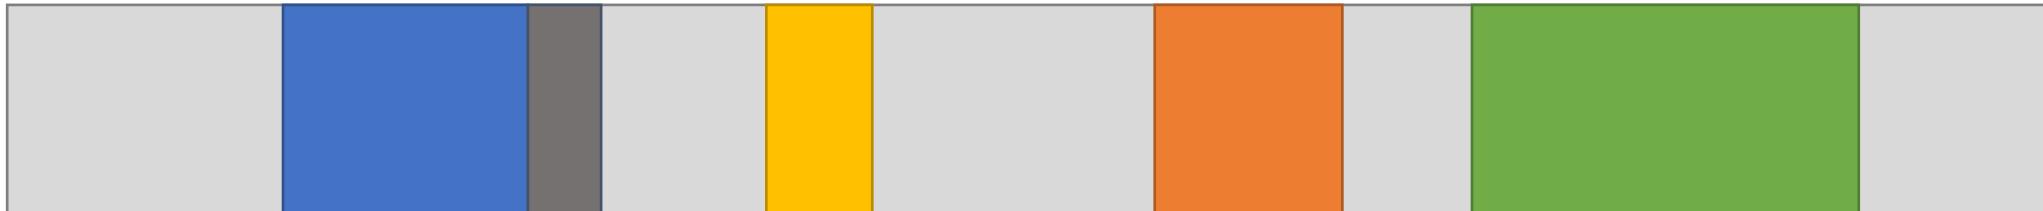
Implementing Compaction (#1): Keep a “forwarding pointer” in each object

1. Compute new locations of objects
2. Update all pointers
3. Move



Implementing Compaction (#2): Keep a table in free space

- Table maps **groups of consecutive objects** to new offsets



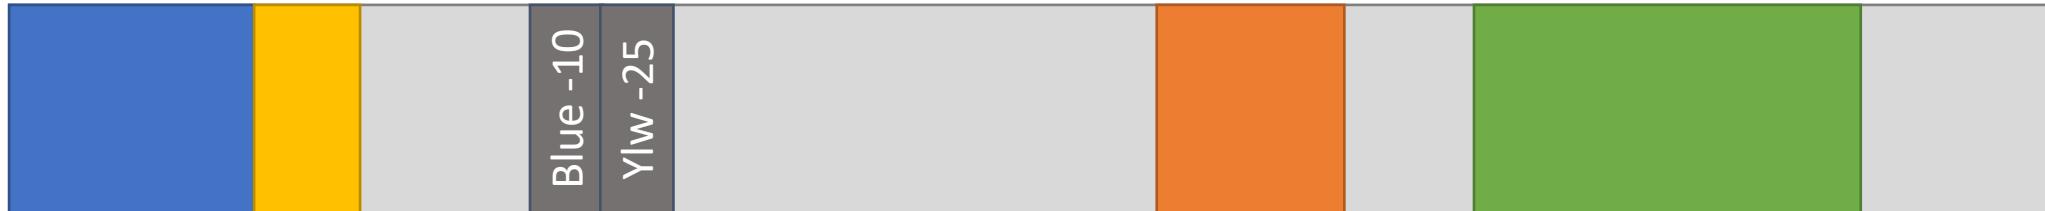
Implementing Compaction (#2): Keep a table in free space

- Table maps **groups of consecutive objects** to new offsets



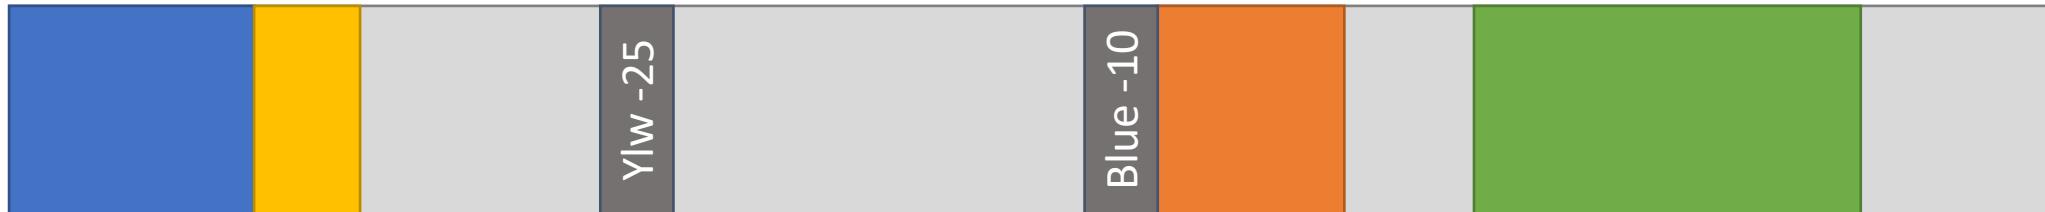
Implementing Compaction (#2): Keep a table in free space

- Table maps **groups of consecutive objects** to new offsets



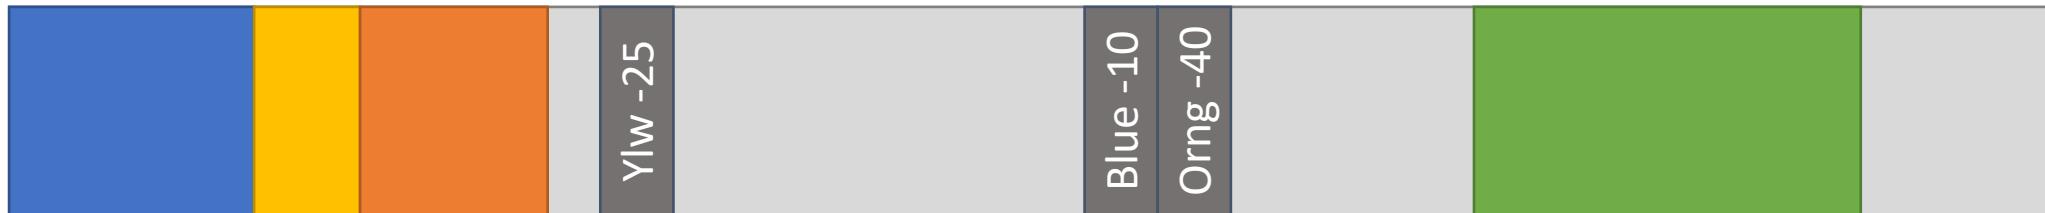
Implementing Compaction (#2): Keep a table in free space

- Table maps **groups of consecutive objects** to new offsets
- “Roll” the table into free space if needed



Implementing Compaction (#2): Keep a table in free space

- Table maps **groups of consecutive objects** to new offsets
- “Roll” the table into free space if needed



Implementing Compaction (#2): Keep a table in free space

- Table maps **groups of consecutive objects** to new offsets
- “Roll” the table into free space if needed
- Need to sort the table at the end



Compacting allows for **really** fast allocation

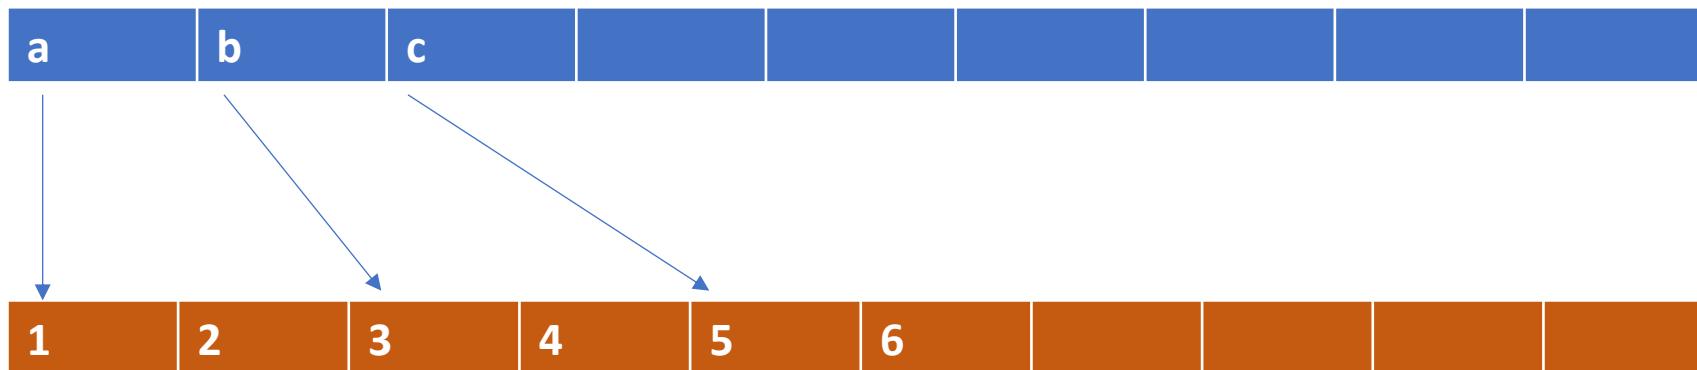
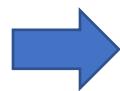
- “Bump allocation”
 - Heap pointer points to end of heap
 - To allocate N bytes:
 - Increment (“bump”) heap pointer by N
 - If we pass the end of the heap, trigger a GC
 - Return old value of heap pointer

... yes. That's it. That's how we implement `malloc`

```
_malloc:  
    lw t0,heapptr          # t0 = heap ptr  
    lw t2,hearend          # t2 = end of heap  
    add t1,t0,a0           # t1 = heap ptr + Nbytes  
    blt t2,t1,__eom       # check if t1 > heap limit  
    sw t1,heapptr          # heap ptr += Nbytes  
    addi a0,t0,0            # a0 = old heap ptr  
    jalr zero,ra,0          # return  
__eom:  
    # trigger GC  
...
```

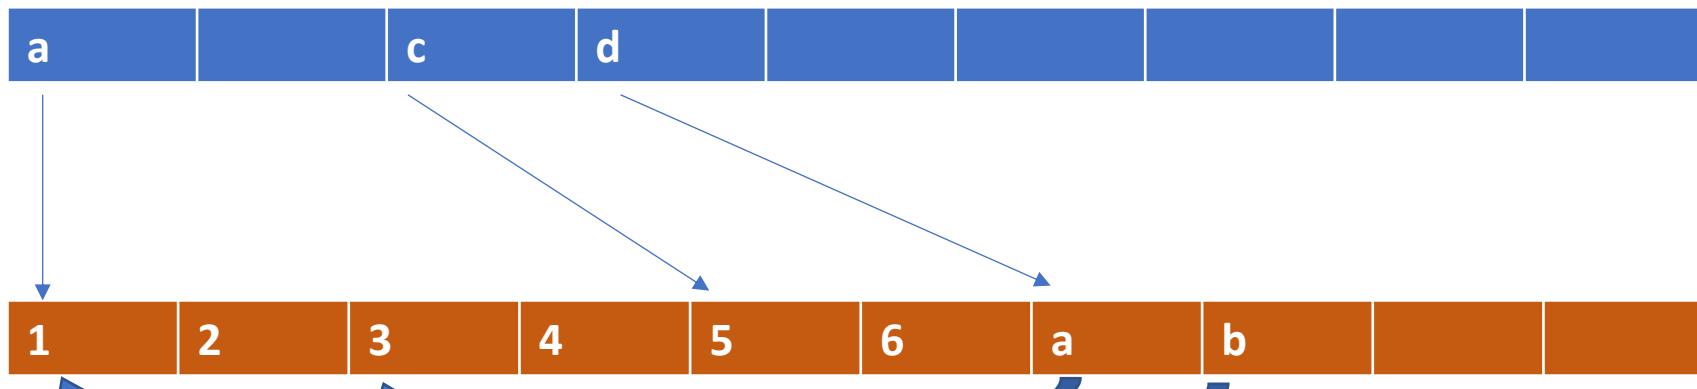
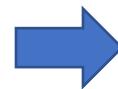
Bump allocation

```
let a = (1, 2)
let b = (3, 4)
let c = (5, 6)
let d = (a, b)
let d = (a, c)
let e = (7, 8)
```



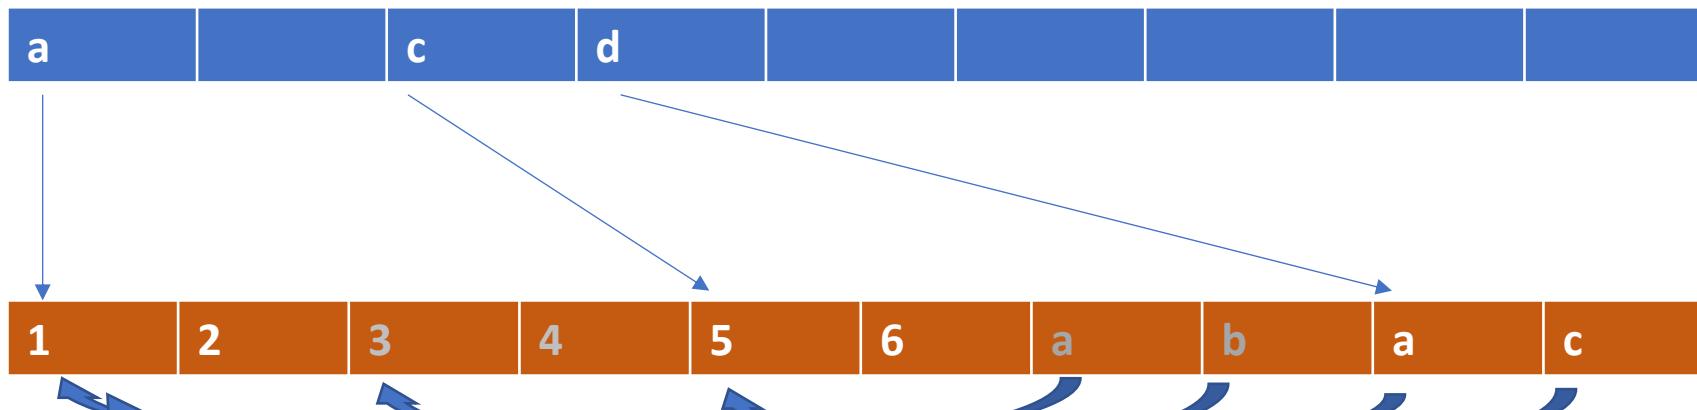
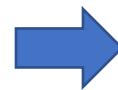
Bump allocation

```
let a = (1, 2)
let b = (3, 4)
let c = (5, 6)
let d = (a, b)
let d = (a, c)
let e = (7, 8)
```



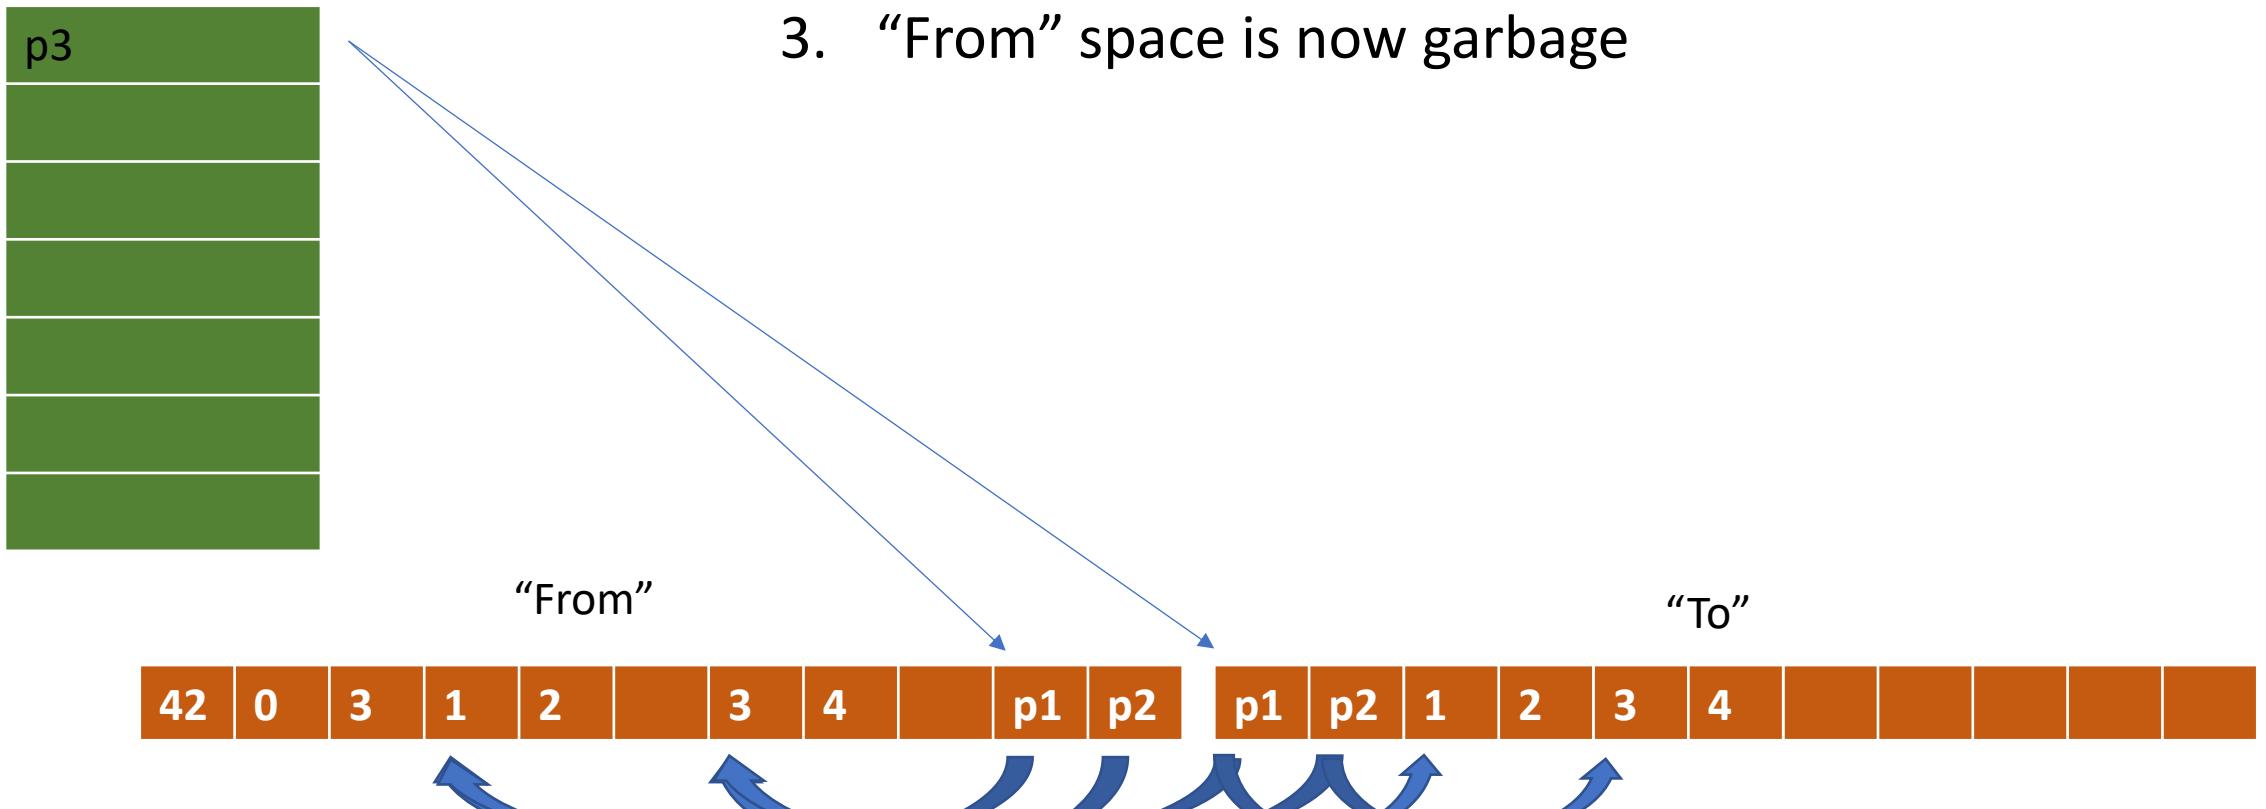
Bump allocation

```
let a = (1, 2)  
let b = (3, 4)  
let c = (5, 6)  
let d = (a, b)  
let d = (a, c)  
let e = (7, 8)
```



GC Strategy #3: Copying

1. Divide heap into “from” space and “to” space
 2. **Copy** live objects into “to” space
 3. “From” space is now garbage



Copying pros and cons

- Pros
 - No traversing of whole heap
 - No fragmentation
- Cons
 - Heap size basically cut in half
 - Have to move pointers

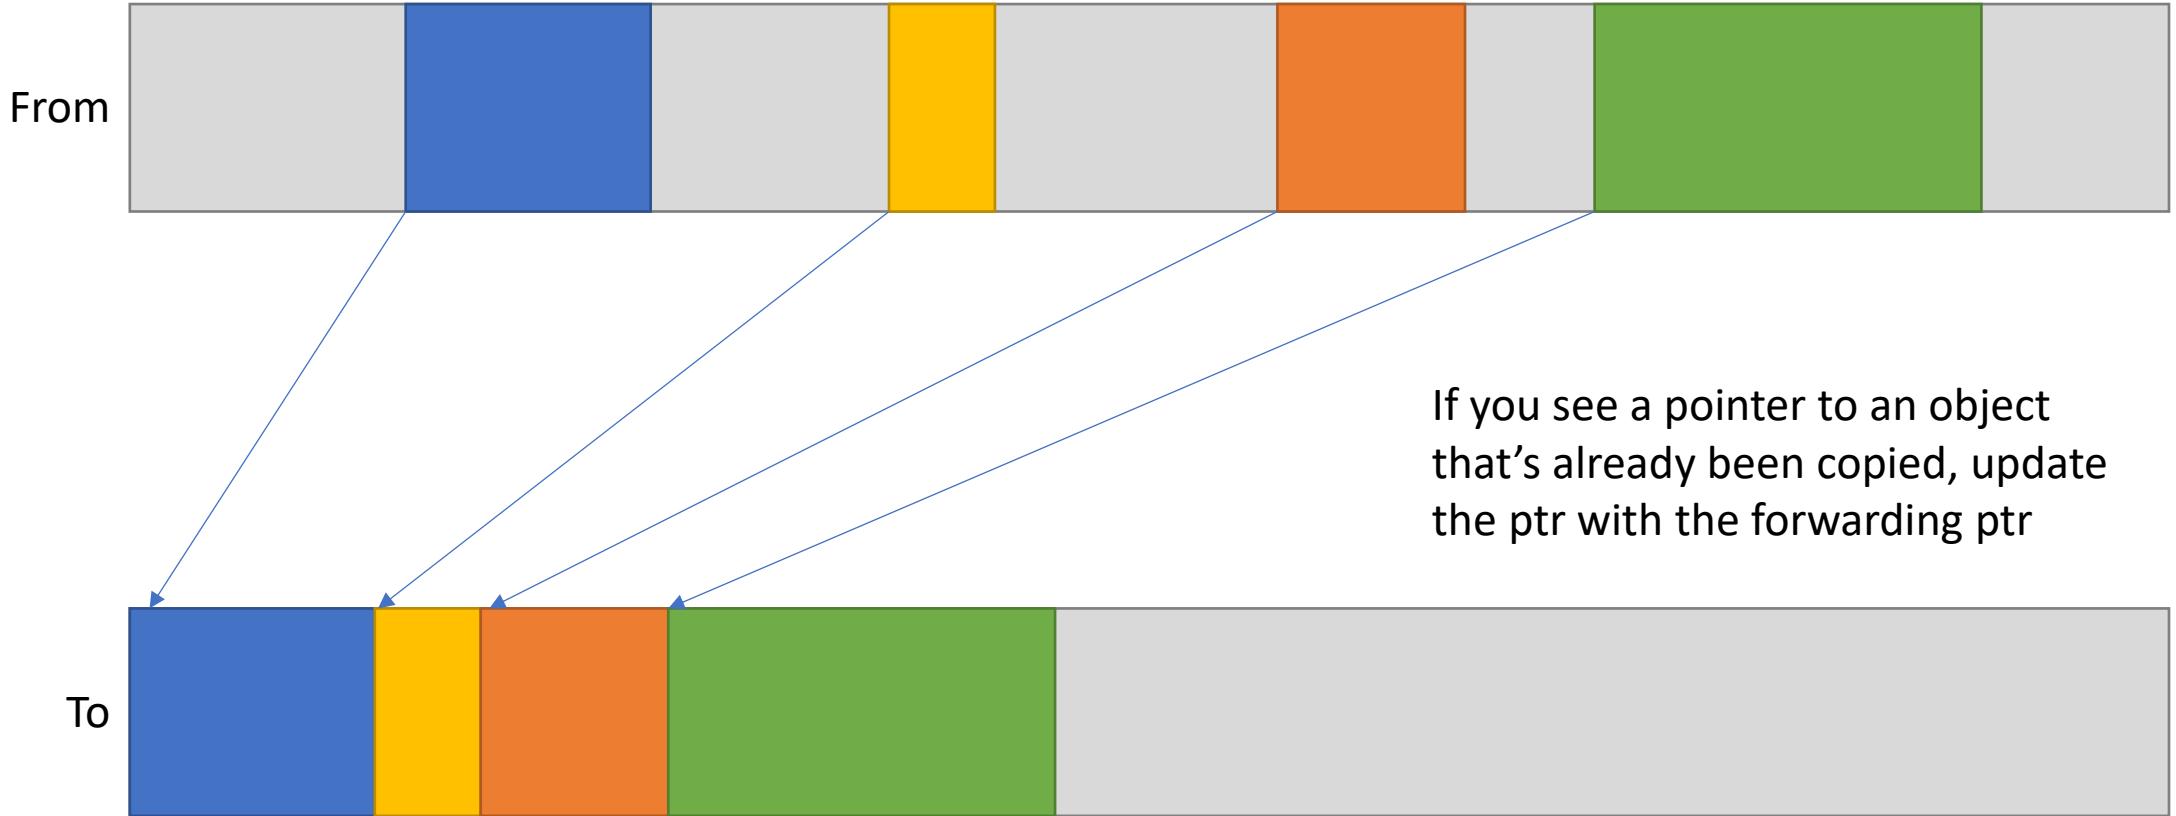
Copying Implementation: Just turn the from space into forwarding pointers



Copying Implementation: Just turn the from space into forwarding pointers



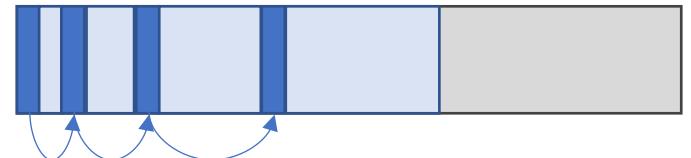
Copying Implementation: Just turn the from space into forwarding pointers



Another side benefit of copying

```
let rec list l n =
  if n <= 0 then l
  else
    list ((List.length
            (List.init (n mod 5) id))::l)
          (n - 1)
```

```
let l = list [] 10000
do_n_times 3 (fun _ -> traverse l)
(* Do other stuff *)
do_n_times 3 (fun _ -> traverse l)
```



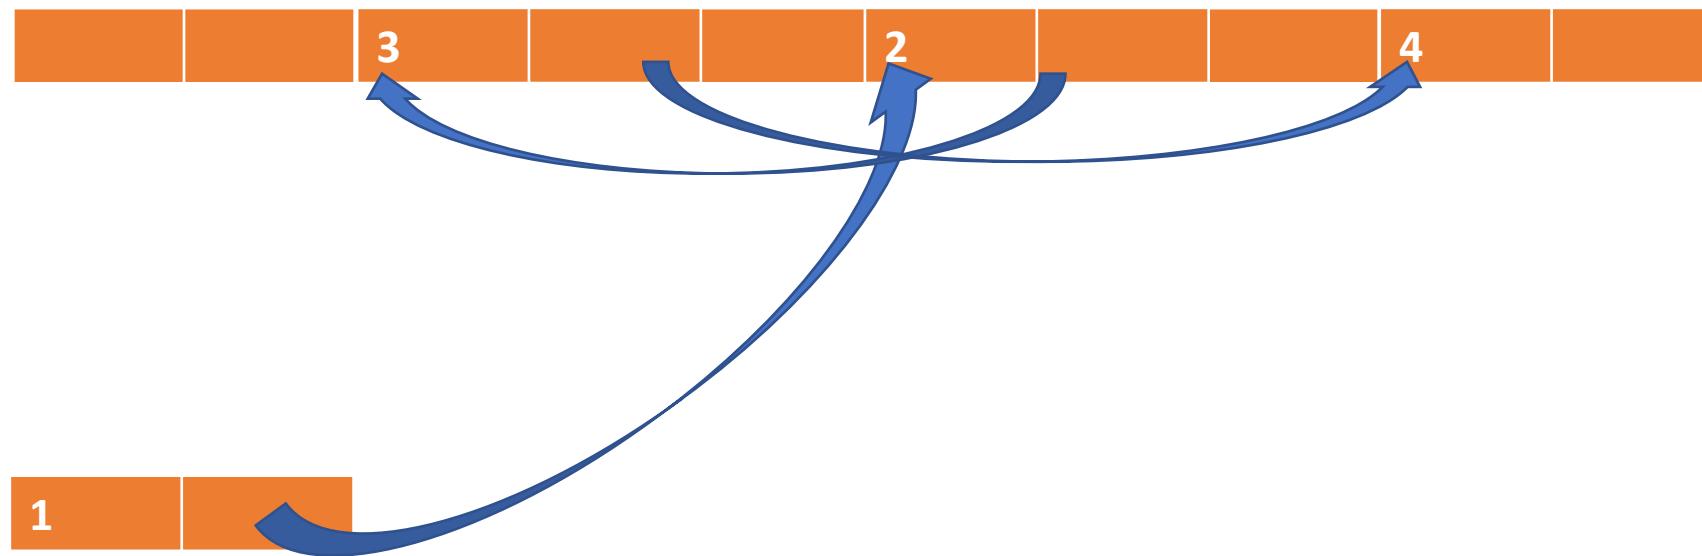
What happened here?

```
Traversed list in 0.00016s
Traversed list in 0.00016s
Traversed list in 0.00016s
Starting new major GC cycle
Traversed list in 0.00007s
Traversed list in 0.00006s
Traversed list in 0.00006s
```

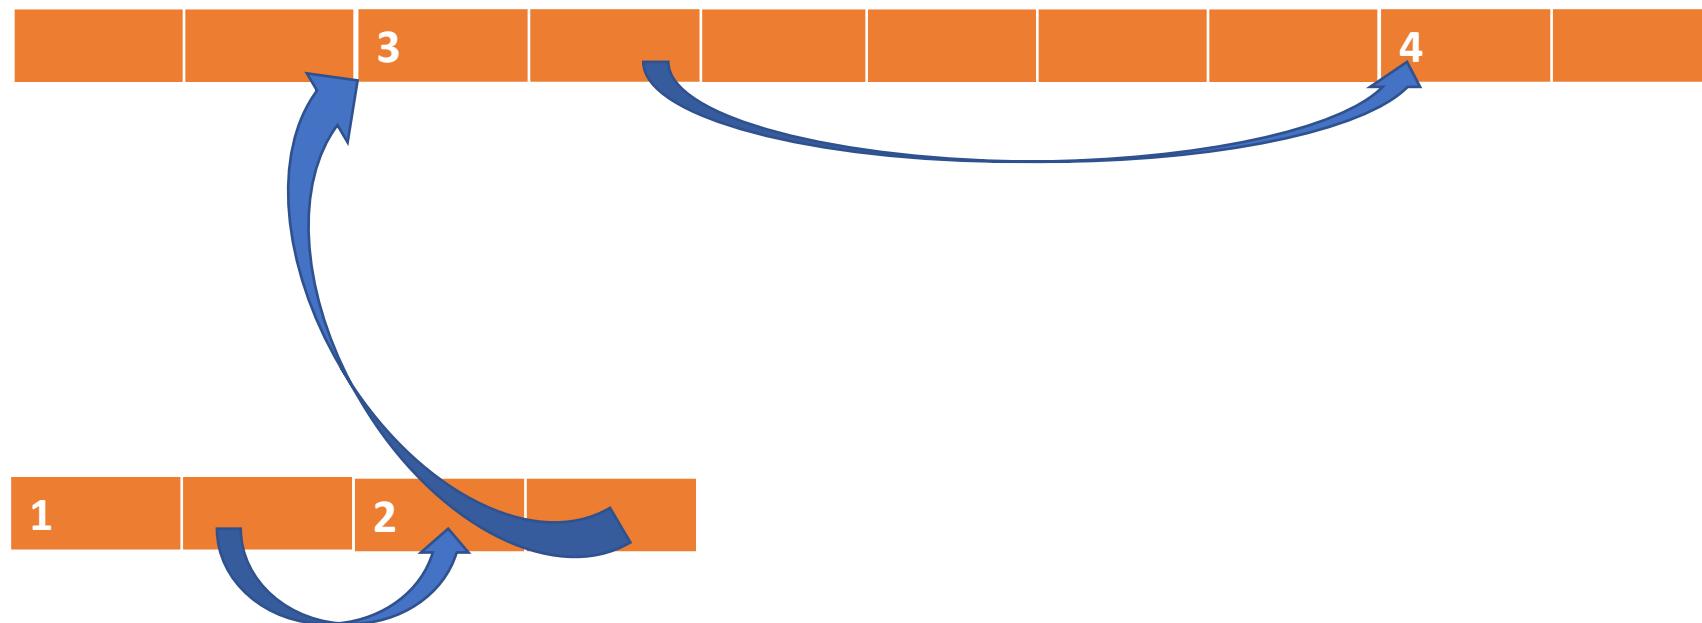
Another side benefit of copying



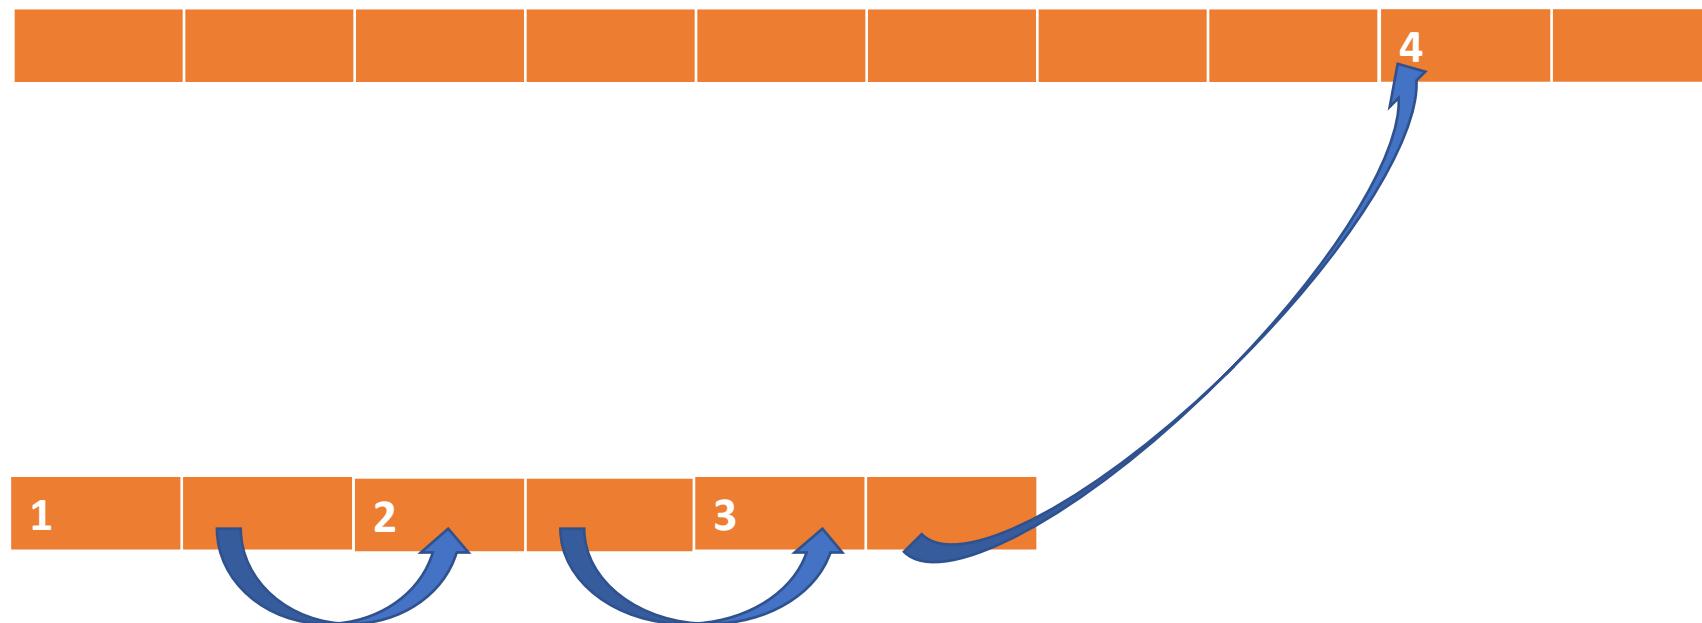
Another side benefit of copying



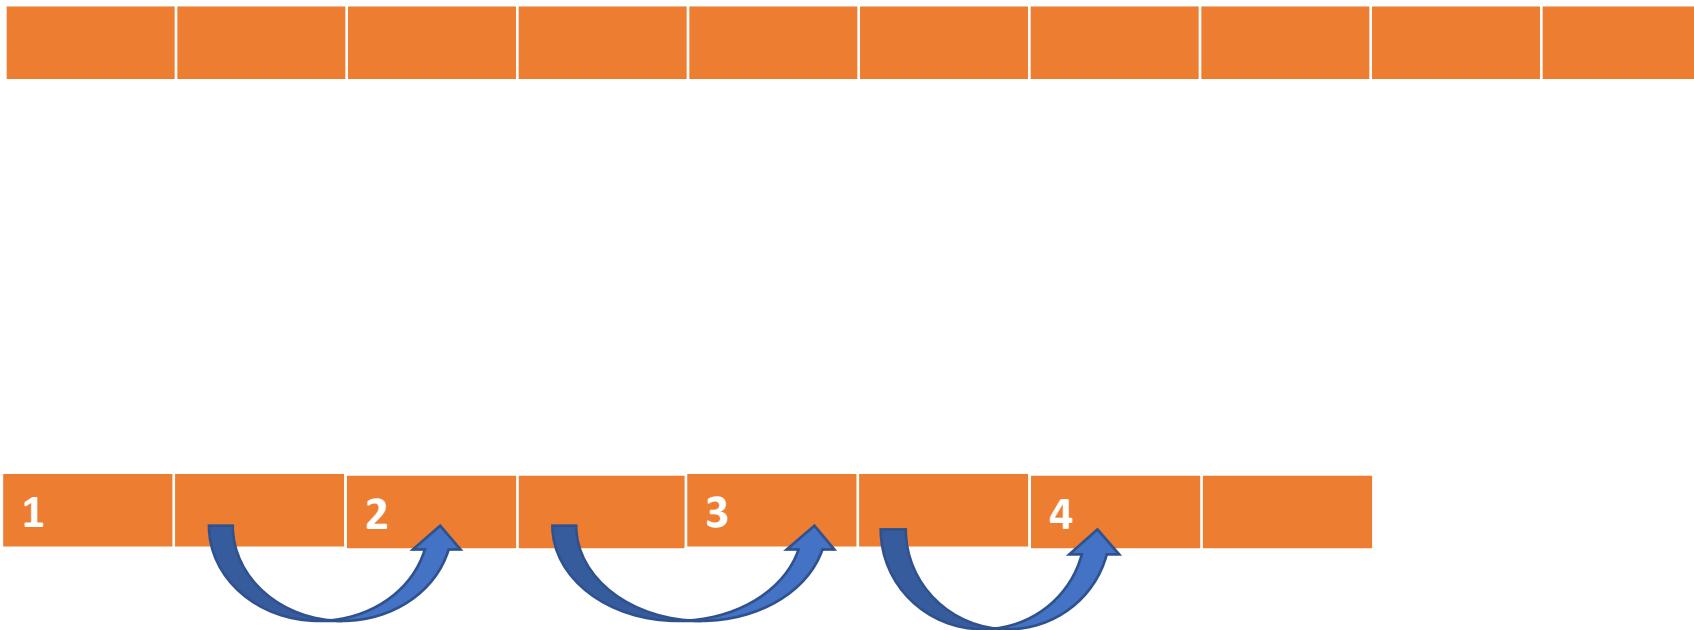
Another side benefit of copying



Another side benefit of copying

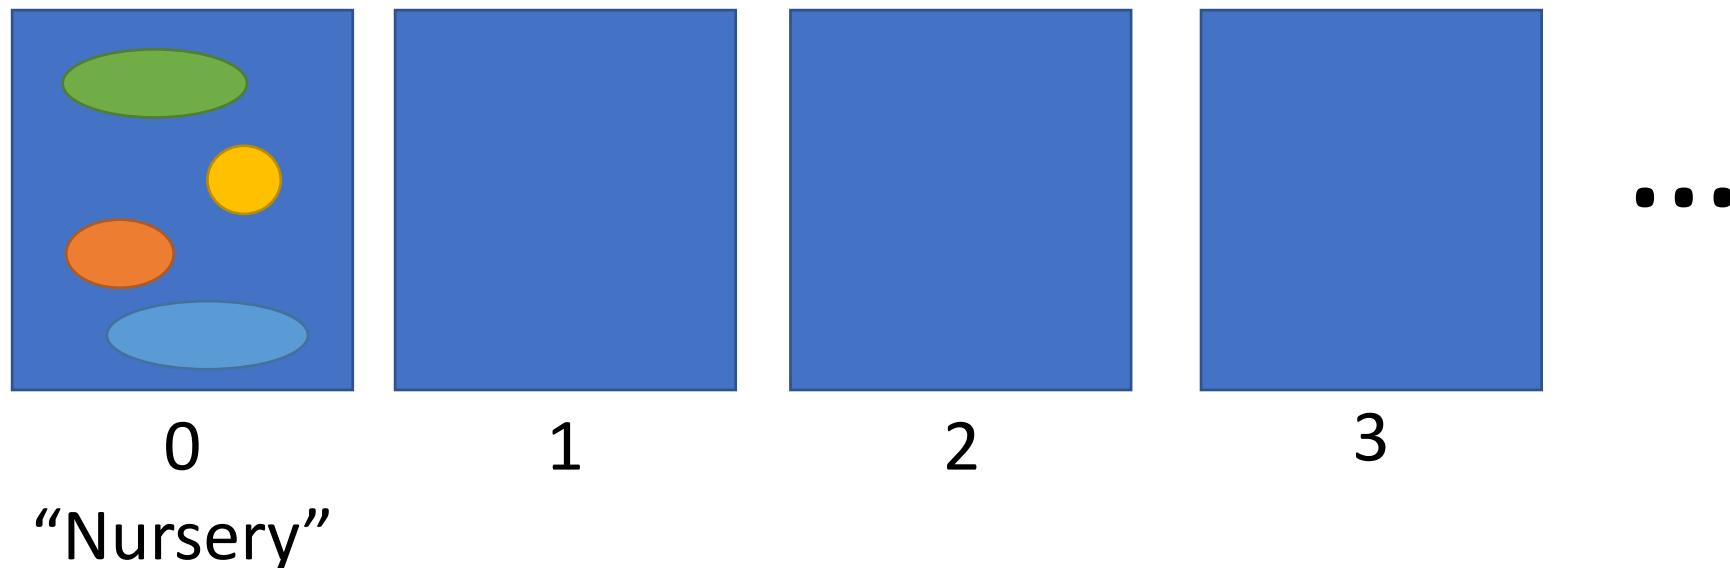


Another side benefit of copying



Generational garbage collection

- Idea: “most objects ‘die young’”
- Separate heap into areas called *generations*
- Collect younger generations more aggressively/frequently



Efficiency

- Most GCs we have discussed are “stop the world”
 - Stop program, do a collection
 - *Pause time*: amount of time a program must wait for the collector
- To reduce pause time, many real-world GCs are *concurrent* or *incremental* (do small amounts of work as the program runs)

In practice, pause times are pretty short

- Don't let people tell you GC makes it totally impractical to use functional languages for real code

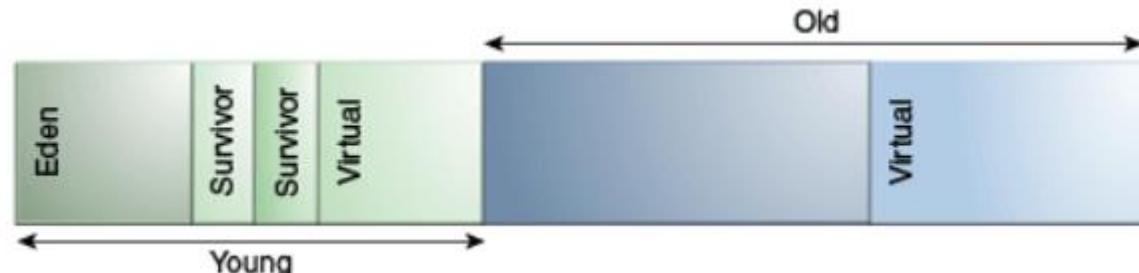
GC type	time ms	number	bytes	bytes/sec
copying	3,063	37	2,111,703,368	689,423,253
mark-compact	0	0	0	-
minor	0	11	4,520	-
total time: 19,902 ms				
total GC time: 3,472 ms (17.4%)				
max pause time: 433 ms				3472 ms / 37 = 93ms avg.
total bytes allocated: 15,794,832,336 bytes				
max bytes live: 140,663,592 bytes				
max heap size: 1,125,367,808 bytes				

OCaml

- Two generations: *minor heap* and *major heap*
 - Allocate large objects directly into major heap
 - “Minor collections” frequent
 - “Major collections” when necessary
- Major collections are (concurrent) mark-compact
 - Not to be confused with *parallel* GC (GC runs on multiple threads to reduce pause time)

Java (HotSpot JVM)

- Generational
 - Eden (nursery)
 - Live objects copied from Eden to one of two “survivor” spaces
 - Copying collection used to copy between survivor spaces
 - After a certain number of copies, moved to “old” generation
- Several different collection strategies available for different applications



Python

- Reference Counting*
 - Generational
 - Cycles?
 - It's complicated