

# Await and Deadlocks

*CS 536: Science of Programming, Spring 2022*

## A. Why?

- Avoiding interference isn't the same as coordinating desirable activities.
- It's common for one thread to wait for another thread to reach a desired state.
- Care needs to be taken to avoid a program that waits with no hope of completing.

## B. Objectives

At the end of this lecture you should know

- The syntax and semantics of the *await* statement.
- How to draw an evaluation diagram for a parallel program that uses *await*.
- How to recognize deadlocked configurations in an evaluation diagram.
- How to list the potential deadlock predicates for a parallel program that uses *await*.

## C. Synchronization

### The Need for Synchronization

- We've looked at parallel programs whose threads avoid bad interactions.
  - They don't interfere because they don't interact (disjoint programs/conditions).
  - They interact but don't interfere (interference-freedom).
- To support good interaction between threads, we often have to have one thread wait for another one. Some examples:
  - Thread 1 should wait until thread 2 is finished executing a certain block of code.
  - Thread 1 has to wait until some buffer is not empty
  - Thread 2 has to wait until some buffer is not full.
- The general problem is that we often want threads to *synchronize*: We want one thread to wait until some other thread makes a condition come true.
- *Example 1:* For a more specific example, in the following program, the calculation of *u* doesn't start until we finish calculating *z*, even though *u* doesn't depend on *z*.

$[x := \dots \parallel y := \dots \parallel z := \dots]; u = f(x, y); v := g(u, z)$

On the other hand, we can't nest parallel programs, so we can't write

$[[[x := \dots \parallel y := \dots]; u = f(x, y) \parallel z := \dots]; v := g(u, z)]$

which would be a natural way to do the calculations of  $u$  and  $z$  in parallel. In some sense, what we'd like is to run something like

$[x := \dots \parallel y := \dots \parallel \text{wait for } x \text{ and } y; u = f(x, y) \parallel z := \dots]; v := g(u, z)$

## D. The Await Statement

- It's time to introduce a new statement, the *await* statement, whose semantics implements the notion of waiting until some condition is true.
  - Busy wait loops like `while !e {skip} {e}` work but are wasteful.
- **Syntax:** `await e then{s}` where  $e$  is a boolean expression and  $s$  is a statement.
  - $s$  isn't allowed to have loops, *await* statements, or atomic regions.
  - *await* statements can only appear in sequential threads of parallel programs. (i.e., in some thread  $s_k$  in an  $[s_1 \parallel s_2 \parallel \dots]$ .)
- An *await* statement is a *conditional atomic region*. Suppose that some thread begins with *await e then{s}*, then
  - We nondeterministically choose between all the available threads, i.e., there's no insistence that we must check the *await* before trying other threads. (See case 1 of Example 2.)
  - If we choose the thread that begins with *await e then {s}*,
    - If  $e$  is true, then immediately jump to  $s$  and execute all of it.
      - The test, jump, and execution of  $s$  are atomic — the combination executes as one step. e.g., with the configuration below, we can't set  $x$  to 1 between looking up the two  $x$ 's to use for calculating  $x+x$ . (See case 2 of Example 2.)
    - If  $e$  is false, we *block*: We wait until  $e$  is true. Instead, we nondeterministically choose between the other threads and execute one. (See case 3 of Example 2.)
- An *await* is similar to an atomic *if-then* statement, but not identical.
  - With `<if e then {s} else {skip}>`, if  $e$  is false, we execute *skip* and complete. (See case 4, Example 2.)
  - With *await e then {s}*, if  $e$  is false, nothing happens until  $e$  becomes true. (See the note with case 3, Example 2.)

*Example 2:*

- (See the discussion above)
  - Case 1:  $<[\text{await } b \text{ then}\{ x := x+x \} \parallel x := 1], \{b = T, x = 2\} >$   
 $\rightarrow <[\text{await } b \text{ then}\{ x := x+x \} \parallel \text{skip}], \{b = T, x = 1\} >.$
  - Case 2:  $<[\text{await } b \text{ then}\{ x := x+x \} \parallel x := 1], \{b = T, x = 2\} > \rightarrow <[\text{skip} \parallel x := 1], \{b = T, x = 4\} >.$   
 (This is the only transition that executes the *await*.)

- Case 3:  $\langle [\text{await } b \text{ then } \{s\} \parallel x := 1], \{b = F\} \rangle \rightarrow \langle [\text{await } b \text{ then } \{s\} \parallel \text{skip}], \{b = F, x = 1\} \rangle$ .  
(The second configuration is blocked, with no other thread available to unblock it.)
- Case 4:  $\langle [\text{if } b \text{ then } \{x := 0\} \text{ else } \{\text{skip}\} \parallel s'], \{b = F\} \rangle \rightarrow \langle [\text{skip} \parallel s'], \{b = F, x = 0\} \rangle$ .
- Example 3: In the introduction, we looked at a situation where we want to wait for some calculations to finish before starting others.

•  $[x := \dots \parallel y := \dots \parallel \text{wait for } x \text{ and } y; u = f(x, y) \parallel z := \dots]; v := g(u, z)$  can be implemented using

```
x_done := F; y_done := F;
[x := ...; x_done := T \parallel y := ...; y_done := T
 \parallel \text{await } x\_done \wedge y\_done \text{ then } \{u := f(x, y)\} \parallel z := \dots];
v := g(u, z);
```

## E. await, wait, if, and $\langle s \rangle$

### The Abbreviations $\langle s \rangle$ and $\text{wait } e$

- With  $\text{await } e \text{ then } \{s\}$ , there are two simple cases: When  $e$  is trivial and when  $s$  is trivial.
- *Definition:* We can redefine  $\langle s \rangle$  to stand for  $\text{await } T \text{ then } \{s\}$ . When the test is trivial, we don't need to wait, we simply execute the body atomically. So atomic execution is just conditional atomic execution with a trivial test.
- *Definition:*  $\text{wait } e \equiv \text{await } e \text{ then } \{\text{skip}\}$ . When the body is trivial, we simply wait; when  $e$  is true, execution is complete.
- There's an important difference between  $\text{wait } e; s$  and  $\text{await } e \text{ then } \{s\}$  (and even  $\text{wait } e; \langle s \rangle$ ).
  - With  $\text{await } e \text{ then } \{s\}$ , once  $e$  is true, we immediately atomically execute  $s$ , so no other statement can interleave between the test and running  $s$ . Therefore  $s$  can rely on  $e$  being true when it starts executing.
  - $\text{wait } e; s$  means  $\text{await } e \text{ then } \{\text{skip}\}; s$ , so it allows another thread to be executed after the  $\text{wait}$  but before running  $s$ .
  - This is still the case with  $\text{wait } e; \langle s \rangle$ . We can't interleave anything with  $s$ , but we can execute another thread between the test and  $s$ .
  - This is a common source of bugs in real parallel programs!

## F. Await Statement Proof Rule and Outlines

- The proof rule for the *await* statement is similar to an *if*, but there's no false clause.

*Await Statement (a.k.a. Synchronization Rule)*

$$\frac{\{p \wedge e\} s \{q\}}{\{p\} \text{await } e \text{ then } \{s\} \{q\}} (\text{Await})$$

- Minimal Proof Outline:*  $\{p\} \text{await } e \text{ then } \{s\} \{q\}$
- Full Proof Outline:*  $\{p\} \text{await } e \text{ then } \{\{p \wedge e\} s^* \{q\}\} \{q\}$  where  $s^*$  is a full proof outline for  $s$ .
- Weakest Preconditions:*  $wp(\text{await } e \text{ then } \{s\}, q) \equiv e \rightarrow wp(s, q)$ .
  - This guarantees  $\{e \rightarrow wp(s, q)\} \text{await } e \text{ then } \{\{wp(s, q)\} s^* \{q\}\} \{q\}$

## G. The Producer/Consumer Problem

- The *Producer/Consumer Problem* (a.k.a. *Bounded Buffer Problem*) is a standard problem in parallel programming.
- We have two threads running in parallel: The producer creates things and puts them into a buffer; the consumer removes things from the buffer and does something with them.
- The problem is that if the buffer is full, the producer shouldn't add anything to the buffer; if the buffer is empty, the consumer shouldn't remove anything from the buffer.
- Example 5:* The rough code to solve this problem is

```
Initialize(buffer);
[while ¬done {                                // Producer
    created := Create();
    await NotFull(buffer) then {
        BufferAdd(buffer, created)
    }
}
|| while ¬done {                                // Consumer
    await NotEmpty(buffer) then
        removed := BufferRemove(buffer);
    };
    Consume(removed)
}
]
```

- Buffer operations need to be synchronized because the threads share the buffer. The threads don't share the created or removed objects, so the *Create* and *Consume* calls can go outside the *await* and interleave execution.

## H. Deadlock

### Blocked Threads; Deadlock

- Recall that  $\langle [await e \text{ then } \{s\}; \dots \| \dots], \sigma \rangle$  is blocked (must wait) if  $\sigma(e) = F$ .
  - If some other thread can make  $e$  true, then the await may eventually unblock.
  - E.g.,  $\langle [await x > y \text{ then } \{s\}; \dots \| \dots; x := y+1; \dots], \sigma \rangle$  could unblock.
  - But if all the other threads have either completed or are themselves blocked, then there's no way for our *await* to unblock. e.g.,  $\langle [await e \text{ then } \{s\} \text{ end}; \dots \| E], \sigma \rangle$  can't evaluate further.
- Definition:* A parallel program is *deadlocked* if it has not finished execution and there's no possible evaluation step to take, i.e., all the threads are either complete or blocked and at least one thread is blocked.
- Example 6:* If  $A \equiv \text{await } x \geq 0 \dots$ , then there's no arrow out of  $\langle [A \| A], \sigma[x \mapsto -1] \rangle$ , so this configuration is deadlocked. If the value of  $x$  had been  $\geq 0$ , then both *await* statements would have been eligible for execution.  
Since only one blocked thread is required for deadlock,  $\langle [await x \geq 0 \dots \| \text{skip}], \sigma[x \mapsto -1] \rangle$  is also deadlocked.

- Threads can block themselves (trivial example: *await false then {s}*).
- More often, threads block because they're waiting for conditions they expect other threads to establish, e.g., if we're running in a state where  $y = 0$  and  $x = 0$ , then these two threads deadlock:
  - Thread 1:  $\{p_1\} \text{ await } y \neq 0 \text{ then } \{x := 1\}; \dots$
  - Thread 2:  $\{p_2\} \text{ await } x \neq 0 \text{ then } \{y := 1\}; \dots$
- A program might deadlock under all execution paths or only certain execution paths.

- Example 7:* The program


$[\text{await } y \neq 0 \text{ then } \{x := 1\} \| \text{await } x \neq 0 \text{ then } \{y := 1\}]$

deadlocks iff you execute in a state where  $x$  and  $y$  are both zero.

- Example 8:* If thread 1 sets  $x := 0$  before thread 2 evaluates its *wait x*, then thread 2 will block. (Recall  $\text{wait } x \equiv \text{await } x \text{ then } \{\text{skip}\}$ .)

```
{T} x := 1; y := 1;
[wait y=1; x:=0 \| wait x=1; y:=0]
{x=0 \wedge y=0}
```

Figure 2 contains an execution graph for this program in state  $\{x=1, y=1\}$  (somewhat abbreviated). There are two deadlocking paths (and four paths that terminate correctly).



- Deadlock is another way for a program to fail to terminate, so  
 $M([wait y=1; x:=0 \parallel wait x=1; y:=0], \{x=1, y=1\}) = \{\{x=1, y=1\}, \perp_d\}$
- Obviously, we'd like to know if a program is going to deadlock. The following test identifies a set of predicates that indicate potential problems with a program; if none of these predicates is satisfiable, then deadlock is guaranteed not to occur.
- If one or more of these predicates is satisfiable, then we can't guarantee that deadlock will not occur, but we aren't guaranteeing that deadlock *must* occur. (So the deadlock conditions are sufficient to show deadlock is impossible but they are not necessary conditions.)
- Let  $\{p\} [ \{p_1\} s_1 * \{q_1\} \parallel \{p_2\} s_2 * \{q_2\} \parallel \dots \parallel \{p_n\} s_n * \{q_n\} ] \{q\}$  be a full outline for a parallel program, where  $p \equiv p_1 \wedge \dots \wedge p_n$  and  $q \equiv q_1 \wedge \dots \wedge q_n$ .
- *Definition:* A (potential) deadlock condition for the program outline above is a predicate of the form  $r'_1 \wedge r'_2 \wedge \dots \wedge r'_n$  where each  $r'_k$  is either
  - $q_k$ , the postcondition for thread  $s_k$  or
  - $p \wedge \neg e$  where  $\{p\} \text{ await } e \dots$  appears in the proof outline for thread  $s_k$ .
  - In addition, at least one of the  $r'_k$  must involve waiting, i.e.,  $q \equiv q_1 \wedge \dots \wedge q_n$  is not a potential deadlock condition.
- A program outline is *deadlock-free* if every one of its potential deadlock conditions is unsatisfiable (i.e., a contradiction):
  - i.e., for each deadlock condition  $r'$ , we have  $\models \neg r'$  (or the equivalent  $\models r' \rightarrow F$ ).

## Parallelism with Deadlock Freedom

$$\frac{[p_i]s_i^*[q_i] \quad [p_i]s_i^*[q_i] \text{ are pairwise interference-free and program is deadlock-free}}{[p_1 \wedge p_2 \wedge \dots \wedge p_n] \ [s_1 \parallel \dots \parallel s_n] \ [q_1 \wedge q_2 \wedge \dots \wedge q_n]}$$

### I. Examples of Deadlock Conditions

- Example 9: Let's take the program from Example 7:

`[await y ≠ 0 then {x := 1} || await x ≠ 0 then {y := 1}]`

and develop an annotation for it:

```
{T}
[ {T} await y ≠ 0 then {{y ≠ 0}} x := 1 {x ≠ 0 ∧ y ≠ 0} ] {x ≠ 0 ∧ y ≠ 0}
|| {T} await x ≠ 0 then {{x ≠ 0}} y := 1 {x ≠ 0 ∧ y ≠ 0} ] {x ≠ 0 ∧ y ≠ 0}
```

- Let set  $D_1 = \{x \neq 0 \wedge y \neq 0, y=0\}$  be the choices for  $p_1'$ .
  - $x \neq 0 \wedge y \neq 0$  is the thread postcondition
  - $y=0$  indicates thread 1 is blocked at the `await` statement.
- Similarly, let set  $D_2 = \{x \neq 0 \wedge y \neq 0, x=0\}$  be the choices for  $p_2'$  (the postcondition of thread 2 and the blocking condition for its `await`).
- There are three choices for the potential deadlock predicate  $r_1' \wedge r_2'$ :
  - $(x \neq 0 \wedge y \neq 0) \wedge (x=0)$ , which is a contradiction.
  - $(y=0) \wedge (x \neq 0 \wedge y \neq 0)$ , which is a contradiction.
  - $(y=0) \wedge (x=0)$ , which is not a contradiction, therefore, it's a potential deadlock condition, and our program does not pass the deadlock-freedom test.
- Recall  $(x \neq 0 \wedge y \neq 0) \wedge (x \neq 0 \wedge y \neq 0)$  is not a potential deadlock predicate because it says that the two threads have both completed.
- One way out of this predicament is to make the initial precondition the negation of  $y=0 \wedge x=0$ . Let  $p$  be  $(x \neq 0 \vee y \neq 0)$  in
 

```
{p}
[ {p} await y ≠ 0 then {p ∧ y ≠ 0} x := 1 {x ≠ 0 ∧ y ≠ 0} end {x ≠ 0 ∧ y ≠ 0}
|| {p} await x ≠ 0 then {p ∧ x ≠ 0} y := 1 {x ≠ 0 ∧ y ≠ 0} end {x ≠ 0 ∧ y ≠ 0}
] {x ≠ 0 ∧ y ≠ 0}
```
- Let  $D_1 = \{x \neq 0 \wedge y \neq 0, p \wedge y=0\}$  and let  $D_2 = \{x \neq 0 \wedge y \neq 0, p \wedge x=0\}$ .
  - The three potential deadlock predicates are now contradictory
    - $(x \neq 0 \wedge y \neq 0) \wedge (p \wedge x=0)$  (is false because of  $x \neq 0 \wedge x=0$ )
    - $(p \wedge y=0) \wedge (x \neq 0 \wedge y \neq 0)$  (is false because of  $y=0 \wedge y \neq 0$ )

- $(p \wedge y=0) \wedge (p \wedge x=0)$   
 $\equiv ((x \neq 0 \vee y \neq 0) \wedge y=0) \wedge ((x \neq 0 \vee y \neq 0) \wedge x=0)$   
 $\Rightarrow (x \neq 0 \wedge y=0) \wedge (y \neq 0 \wedge x=0)$   
 $\Rightarrow F$
- (end of example 9)
- *Example 10:* Since it has three threads, the deadlock conditions for this program are a bit more involved than for Example 9. Thread 1 has one `await` statement, thread 2 has two `await` statements, and thread 3 has no `await` statements.
 

```
[ ... {p11} await e11 ... {q1}
    || ... {p21} await e21 ... {p22} await e22 ... {q2}
    || ... {q3} ]
```
- The deadlock conditions are built using the three sets
  - $D_1 = \{p_{11} \wedge \neg e_{11}, q_1\}$
  - $D_2 = \{p_{21} \wedge \neg e_{21}, p_{22} \wedge \neg e_{22}, q_2\}$
  - $D_3 = \{q_3\}$ .
- Let  $D$  be the set of deadlock conditions,  $D = \{r_1 \wedge r_2 \wedge r_3 \mid r_1 \in D_1, r_2 \in D_2, r_3 \in D_3\} - \{q_1 \wedge q_2 \wedge q_3\}$ . Specifically, we get the following ( $2 \times 3 \times 1 - 1 = 5$ ) conditions:
 

$D = \{(p_{11} \wedge \neg B_{11}) \wedge (p_{21} \wedge \neg B_{21}) \wedge q_3,$	Thread 1 blocked; thread 2 blocked at 1st await
$(p_{11} \wedge \neg B_{11}) \wedge (p_{22} \wedge \neg B_{22}) \wedge q_3,$	Thread 1 blocked; thread 2 blocked at 2nd await
$(p_{11} \wedge \neg B_{11}) \wedge q_2 \wedge q_3,$	Thread 1 blocked
$q_1 \wedge (p_{21} \wedge \neg B_{21}) \wedge q_3,$	Thread 2 blocked at 1st await
$q_1 \wedge (p_{22} \wedge \neg B_{22}) \wedge q_3 \}$	Thread 2 blocked at 2nd await
- The program will be deadlock-free if every predicate in  $D$  is a contradiction (i.e., unsatisfiable).