

Separation Logic Introduction

Stefan Muller

(Based on material by John Reynolds and Peter O'Hearn)

CS 536: Science of Programming, Spring 2022
Lecture 26

1 Pointers and Aliasing

So far, we've only dealt with "local variables": we have a state σ that maps variables to values. In a program, these correspond to variables stored on the stack or in registers. Programming languages like C also allow you to access values on the *heap*, through a "pointer" to a location in memory.

We'll add this to our IMP language. We now have a type of "locations", which represents a pointer to a location in the heap (which contains a value, which may be an integer, a Boolean, or another location). We can "dereference" locations with the syntax $[e]$ (this corresponds to `*e` in C).

This leads to a lot of new problems. Consider the following triple. The predicates don't use exactly the syntax we'll end up using, but it'll work for now.

$$\not\models \{[y] = 0\} [x] := 5 \ {[y] = 0}$$

It seems fairly obvious that this triple should be valid, and yet it isn't! The reason is something we talked about briefly when discussing array assignments: *aliasing*. It's the same reason this triple isn't valid:

$$\not\models \{a[i] = 0\} a[j] := 5 \ {a[i] = 0}$$

At runtime, x and y may point to the same location in the heap, so setting $[x]$ to 5 would also set $[y]$ to 5.

With arrays, we solved this problem by adding conditions to the predicates about whether i and j were equal at runtime. Something like that would likely work here too, but researchers in the early 2000s developed a more powerful extension of Hoare logic that's designed to handle exactly this by adding explicit assumptions about (non)-aliasing: *separation logic*.

We'll make the following additions to our language:

$$s ::= \dots \mid [e] := e \mid x := [e] \mid \text{alloc } e \mid \text{free } e$$

The statement $[e_1] := e_2$ evaluates e_1 and e_2 and sets the heap location pointed to by e_1 to e_2 . The statement $x := [e]$ sets x (which is still a normal variable in the state, as before) to the value in the heap pointed to by e . The statement `alloc e` allocates a new location in the heap initialized to the value of e , and `free e` evaluates e to a location and frees it; these are like `malloc` and `free` in C.

The state now has two parts: a *store*, which we'll continue to denote σ , and a heap, which we'll denote h . This impacts our whole operational semantics: the small-step judgment is now $\langle s, (\sigma, h) \rangle \rightarrow \langle s', (\sigma', h') \rangle$ and the big-step judgment is $M(s, \sigma, h)$. Even evaluation of expressions needs to use both the store and heap: $(\sigma, h)(e)$. Finally, we need to consider satisfaction of triples and predicates under both a store and a heap:

$$\begin{aligned} \sigma, h &\models p \\ \sigma, h &\models \{p\} s \{q\} \end{aligned}$$

1.1 Semantics

The semantics of the new statements use and update the heap. We'll use the notation ℓ to represent heap locations.


$$\frac{(\sigma, h)(e) = \ell \quad h(\ell) = e'}{\langle x := [e], (\sigma, h) \rangle \rightarrow \langle \text{skip}, (\sigma[x \mapsto e'], h) \rangle} \quad \frac{(\sigma, h)(e_1) = \ell \quad \ell \in \text{Dom}(h)}{\langle [e_1] := e_2, (\sigma, h) \rangle \rightarrow \langle \text{skip}, (\sigma, h[\ell \mapsto (\sigma, h)(e_2)]) \rangle}$$

$$\frac{\ell \text{ fresh}}{\langle x := \text{alloc } e, (\sigma, h) \rangle \rightarrow \langle \text{skip}, (\sigma[x \mapsto \ell], h[\ell \mapsto (\sigma, h)(e)]) \rangle} \quad \frac{(\sigma, h)(e) = \ell \quad h' = h \setminus \ell}{\langle \text{free } e, (\sigma, h) \rangle \rightarrow \langle \text{skip}, (\sigma, h') \rangle}$$


As an example, let $s = x := \text{alloc } 42; y := x; [y] := 32$ and $\sigma = \{\}$ and $h = \{\}$.

$$\begin{aligned} & \langle s, (\sigma, h) \rangle \\ & \xrightarrow{2} \langle y := x; [y] := 32, (\{x = \ell\}, \{\ell \mapsto 42\}) \rangle \\ & \xrightarrow{2} \langle [y] := 32, (\{x = \ell, y = \ell\}, \{\ell \mapsto 42\}) \rangle \\ & \xrightarrow{2} \langle \text{skip}, (\{x = \ell, y = \ell\}, \{\ell \mapsto 32\}) \rangle \end{aligned}$$

We usually draw heaps with box-and-arrow diagrams like this:



Heap cells can also contain pointers; this even allows us to create cycles.



2 Separation Logic

2.1 Predicates

We'll introduce new types of predicates we can use in conditions:

- emp means the heap is empty (this generally isn't useful on its own, but will be useful in conjunction with the other predicates).
- $x \mapsto e$ means that the heap consists of *exactly* a location x containing e . Again, requiring exactly one heap location is a big restriction, but this will turn out to be useful.
- $p_1 * p_2$ means that the heap can be divided into two *disjoint* heaps h_1 and h_2 (disjoint meaning they contain different locations) and h_1 satisfies p_1 and h_2 satisfies p_2 .


We'll also use $x \hookrightarrow e$ as shorthand for $x \mapsto e * T$; this says that the heap has a binding of x to e and possibly others (if there are others, we split them into the disjoint heap about which we only assert T , which always holds). Finally, we'll use $x \mapsto _$ to say that x is in the heap, but we don't care what value it points to; we could write this formally as $\exists e. x \mapsto e$.

We can use these together with the usual connectives $\wedge, \vee, \rightarrow, \leftrightarrow, \neg$, with their usual meanings. These interact in interesting ways.


- $\sigma, h \models x \mapsto 5 * y \mapsto 5$ means that h can be divided into disjoint heaps h_1 and h_2 , $\sigma(x)$ is a location in h_1 containing 5 and $\sigma(y)$ is a location in h_2 also containing 5. This means that x and y **cannot be aliases!**



- $\sigma, h \models x \mapsto 5 \wedge y \mapsto 5$ means that h is a heap that has exactly one location x containing 5 and is also a heap with exactly one location y containing 5. This means that x and y **must be aliases!**



- $x \mapsto 5 \wedge y \mapsto 10$ is a contradiction.
- $\sigma, h \models x \hookrightarrow 5 \wedge y \hookrightarrow 5$ means that x and y may or may not be aliases.



2.2 Rules

The rules for allocation and lookup both also perform an assignment to a variable, and thus have a feature that should be familiar from the forward assignment rule: because e can use x , we have to make a new variable x_0 to store the old value of x and substitute x_0 for x in e . Otherwise, the rule for allocation says that if we have an empty heap before, the new heap maps x to e . If e maps to v in the heap, the lookup rule says that x is now bound to v . The update rule is fairly straightforward, but also requires that e is in the heap before the update. It turns out this is actually necessary to enforce partial correctness, not just total correctness, as we'll see. The deallocate rule takes a heap with a binding for e to an empty heap.

Of course, all of these rules require pretty strong conditions. Just like with standard triples, we'll want a way to weaken them; for example, we want to be able to update the value of a location in a heap with more than one location. The rule that allows us to do this is called the Frame Rule: it says that if there's a disjoint part of the heap r that isn't affected by the program and if $\{p\} s \{q\}$, then the same program runs on a heap that satisfies $p * r$ and leaves r unchanged.

$$\begin{array}{c}
 \frac{}{\vdash \{x_0 = x \wedge \text{emp}\} x := \text{alloc } e \{x \mapsto [x_0/x]e\}} \text{ (ALLOC)} \\[10pt]
 \frac{\vdash \{x_0 = x \wedge e \mapsto v\} x := [e] \{x = v \wedge [x_0/x]e \mapsto v\}}{\vdash \{x_0 = x \wedge e \mapsto v\} x := [e] \{x = v \wedge [x_0/x]e \mapsto v\}} \text{ (LOOKUP)} \quad \frac{\vdash \{e \mapsto __ \} [e] := e' \{e \mapsto e'\}}{\vdash \{e \mapsto __ \} [e] := e' \{e \mapsto e'\}} \text{ (UPDATE)} \\[10pt]
 \frac{}{\vdash \{e \mapsto __ \} \text{ free } e \{\text{emp}\}} \text{ (DEALLOCATE)} \quad \frac{\vdash \{p\} s \{q\} \quad \text{FV}(r) \cap \text{Change}(s) = \emptyset}{\vdash \{p * r\} s \{q * r\}} \text{ (FRAME)}
 \end{array}$$

2.3 Examples

Let's take a look back at the program from the beginning and see what goes wrong when we try to prove it:

$$\{y \mapsto 0\} [x] := 5 \{y \mapsto 0\}$$

We can't prove this using the Update rule because the precondition doesn't say anything about x , so we need to add this to the precondition. There are (at least) three ways we can do this, corresponding to the

cases we saw above:

$$\begin{aligned} p_1 &\equiv y \mapsto 0 \wedge x \mapsto _ \\ p_2 &\equiv y \mapsto 0 * x \mapsto _ \\ p_3 &\equiv y \hookleftarrow 0 \wedge x \hookrightarrow _ \end{aligned}$$

All of these are a bit stronger than our original precondition (they need to be): p_1 and p_3 just say that x is in the context; p_2 , on the other hand, requires that x is not an alias of y . Furthermore, p_1 and p_3 still don't match what we need to apply the rule. On the other hand, if we use p_2 , we can prove a modified version of the triple with the Frame Rule:

$$\frac{\vdash \{x \mapsto _\} [x] := 5 \{x \mapsto 5\} \text{ (UPDATE)}}{\vdash \{x \mapsto _ * y \mapsto 0\} [x] := 5 \{x \mapsto 5 * y \mapsto 0\} \text{ (FRAME)}}$$

This makes sense: the program was incorrect because x and y could be aliases; if we assume in the precondition that they're not, then we can prove the program correct.

3 References

These notes were based heavily on the following two resources, which are good places to look for more information:

1. John C. Reynolds, *Separation Logic: A Logic for Shared Mutable Data Structures*. LICS '02
2. Peter W. O'Hearn, "A Primer on Separation Logic (and Automatic Program Verification and Analysis)". *Software Safety and Security; Tools for Analysis and Verification*. NATO Science for Peace and Security Series, vol. 33, pp. 286–318, 2012.